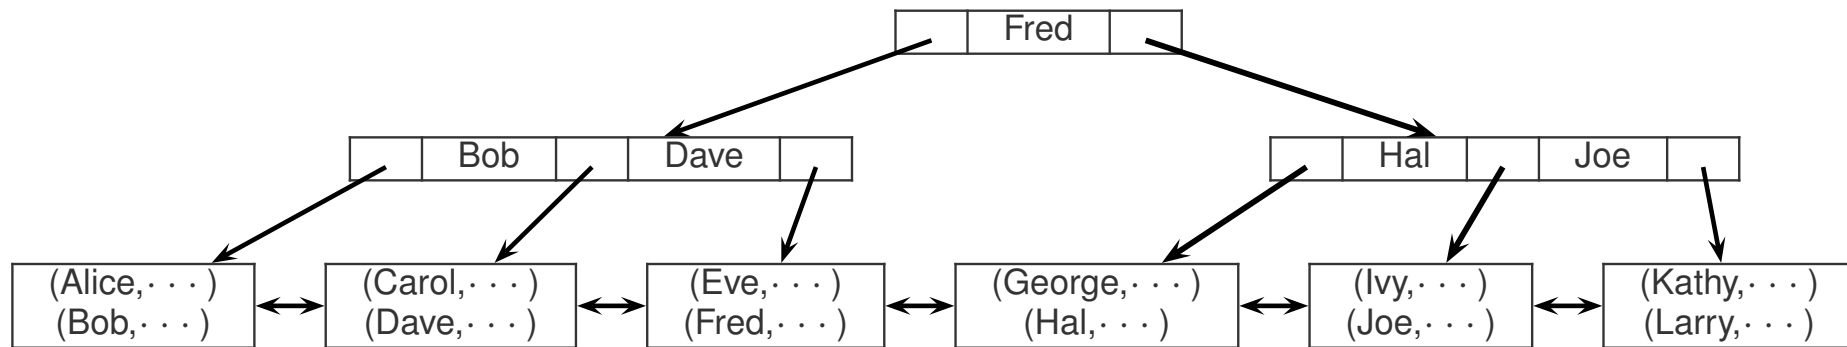


CS4224/CS5424 Lecture 4

Storage & Indexing


B⁺-tree Index



LSM Storage


- LSM = **Log-Structured Merge**
- Inspired by LSM-Tree
 - ▶ P. O’Neil, E. Cheng, D. Gawlick, E. O’Neil, *The Log-Structured Merge-Tree (LSM-Tree)*, Acta Inf., 1996
- Improve write throughput by “converting” random I/O to sequential I/O
 - ▶ Append-only updates instead of in-place updates
- Used in BigTable, Cassandra, DynamoDB, HBase, LevelDB, MyRocks, RocksDB, SQLite4, Voldemort, WiredTiger, YugabyteDB, etc.

LSM-Tree



(O'Neil, Cheng, Gawlick, & O'Neil, 1996)

LSM-Tree (cont.)



(O'Neil, Cheng, Gawlick, & O'Neil, 1996)

LSM Storage

- LSM storage for a relation $R(K, V)$ consists of:
 - ▶ A main-memory structure **MemTable**
 - ▶ A set of disk-based structures **SSTables**
 - ▶ A commit log file
- **MemTable = Memory Table**
 - ▶ Contains the most recent updates organized in main-memory
 - ▶ MemTable is updated in-place
 - ★ Deleted records aren't removed but marked with tombstones (denoted by \perp)
 - ▶ When size of MemTable reaches a certain threshold (e.g., 2MB), the records in MemTable are sorted and flushed to disk as a new SSTable
- A key may have multiple versions of values

SSTable (Sorted String Table)

- SSTables are immutable structures
- SSTable records are sorted by relation's key K
- Each SSTable is associated with a range of key values & a timestamp

Commit Log File

- A **commit log file** is used to ensure durability
- Each new update is appended to commit log & updated to MemTable

LSM Storage: Example

MemTable	SSTable 1	SSTable 2	SSTable 3
<div><div>7, x</div><div>192, ⊥</div></div>	<div><div>5, a</div><div>160, b</div><div>180, d</div></div>	<div><div>160, ⊥</div><div>192, c</div><div>300, a</div></div>	<div><div>7, m</div><div>180, j</div><div>230, n</div></div>

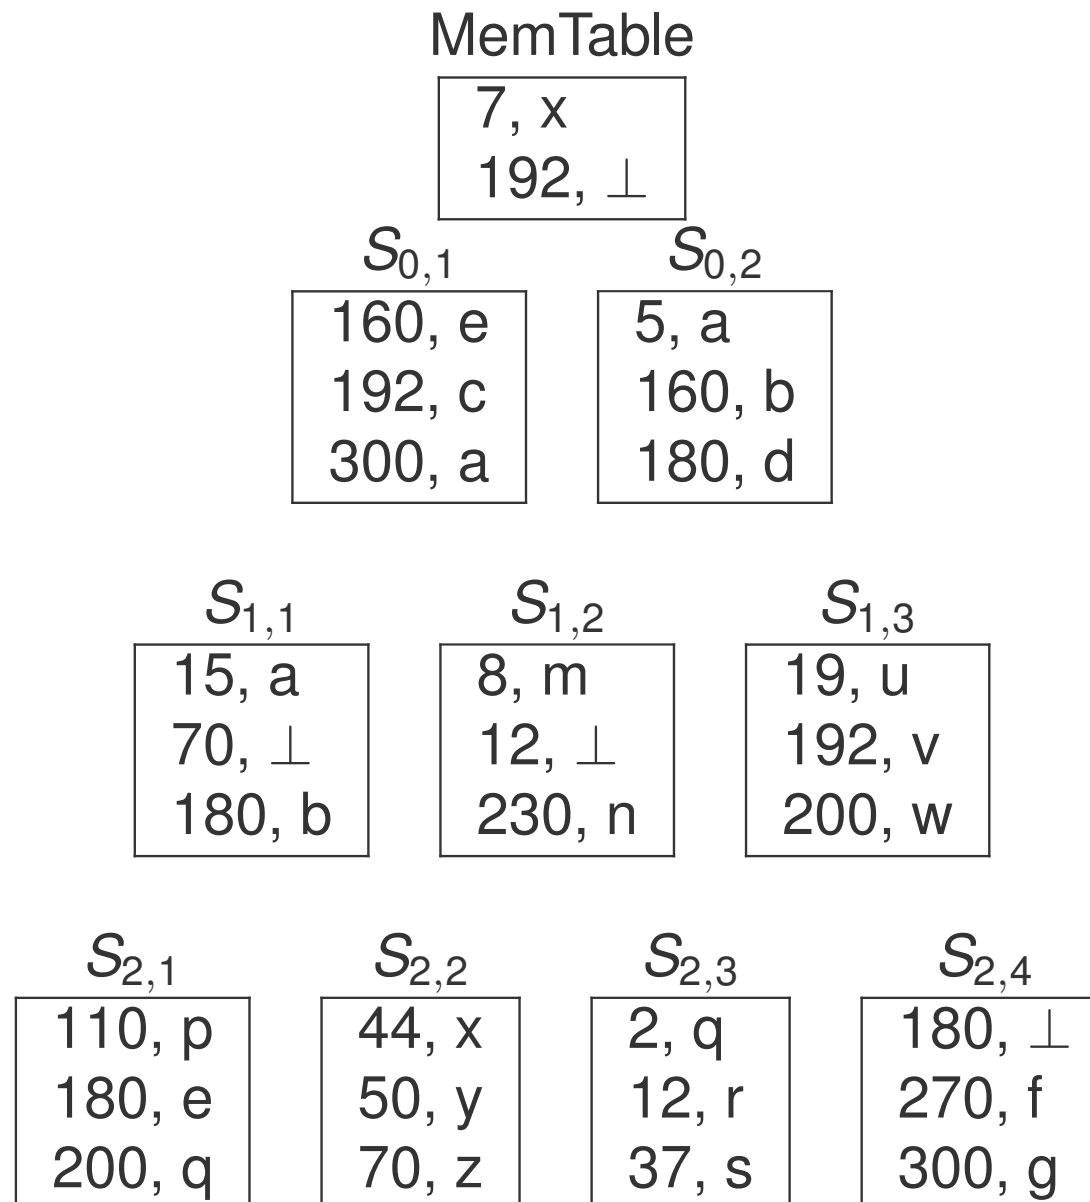
$\text{timestamp}(\text{SSTable 1}) < \text{timestamp}(\text{SSTable 2}) < \text{timestamp}(\text{SSTable 3})$

$\text{Range}(\text{SSTable 1}) = [5, 180]$
 $\text{Range}(\text{SSTable 2}) = [160, 300]$
 $\text{Range}(\text{SSTable 3}) = [7, 230]$

Compaction of SSTables

- Maintenance task to merge SSTable records
 - ▶ Improves read performance by defragmenting table records
 - ▶ Improves space utilization by eliminating tombstones & stale values
- Compaction Strategies
 - ▶ Size-tiered Compaction Strategy (STCS)
 - ▶ Leveled Compaction Strategy (LCS)
 - ▶ etc.



Compaction organizes SSTables into tiers




Size-Tiered Compaction Strategy (STCS)

- SSTables are organized into tiers with SSTables in each tier having approximately the same size
- Compaction is triggered at a tier L when the number of SSTables reaches a threshold (e.g., 4)
 - ▶ All SSTables in tier L are merged into a single SSTable that is stored in tier $L + 1$
 - ▶ Tier L becomes empty after compaction

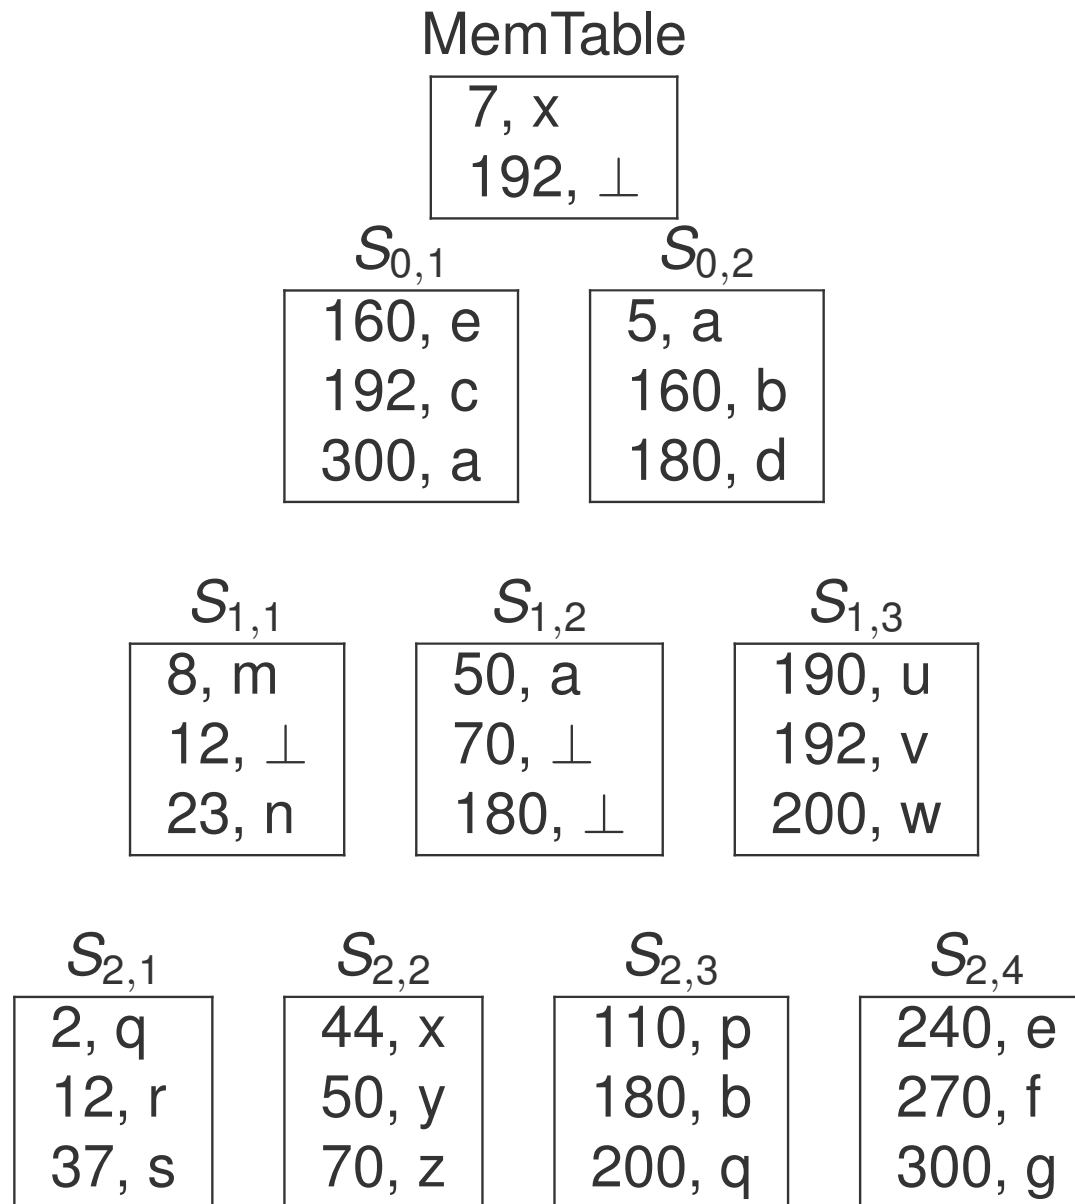
Size-Tiered Compaction: Example



Example: Merging SSTables



Leveled Compaction Strategy (LCS)



Leveled Compaction Strategy (LCS)

- SSTables are organized into a sequence of levels: level-0, level-1, etc.
- Two SSTables **overlap** if their key ranges overlap
- SSTables at level 0 may overlap
- For each level $L \geq 1$
 - ▶ Each SSTable has the same size (e.g., 2MB)
 - ▶ SSTables at the same level do not overlap
 - ▶ Each SSTable at level L overlaps with at most F SSTables at level $L+1$ (F = **compaction factor**)
- If a key appears in two SSTables at different levels i & j , $i < j$, the version at level i is more recent
- $S_{i,j}$ is more recently created than $S_{i,k}$ if $j > k$

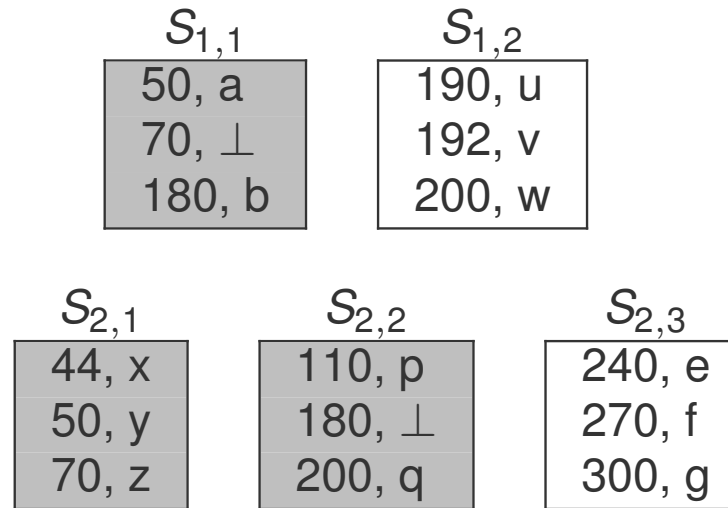
Leveled Compaction of SSTables

- How to perform compaction at level L ?
- $L \geq 1$:
 - ▶ Select a SSTable S at level L
 - ★ Let v be the ending key of the last compaction at level L
 - ★ S is the first level- L SSTable that starts after v if it exists; otherwise, S is the level- L SSTable with smallest start key value
 - ▶ Merge S with all overlapping SSTables at level $L + 1$
- $L = 0$:
 - ▶ Merge all SSTables at level 0 with all overlapping SSTables at level 1
- New SSTables are stored at level $L + 1$
- Old SSTables are removed

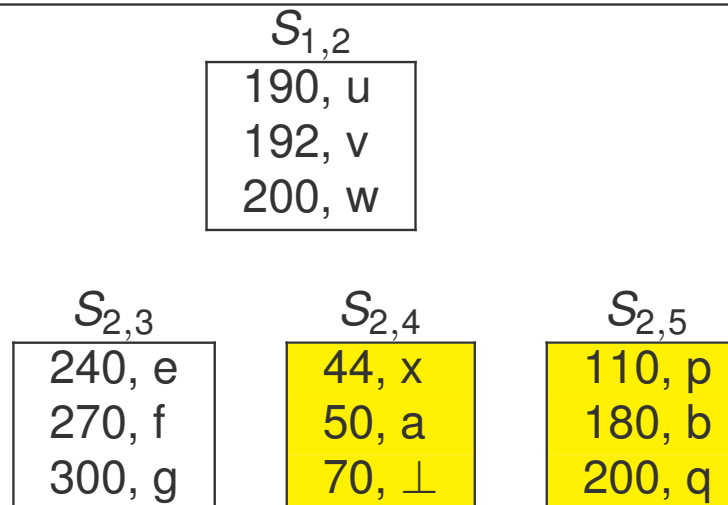
Example: Compaction of $S_{1,1}$

- Merges $S_{1,1}$ with $\{S_{2,1}, S_{2,2}\}$ to $\{S_{2,4}, S_{2,5}\}$

Before
Compaction



After
Compaction



Example: Compaction at Level 0

- Merge all level-0 SSTables with overlapping level-1 SSTables
- **Example:**

Before
Compaction

$\text{Range}(S_{0,1}) = [20, 400]$

$\text{Range}(S_{0,2}) = [12, 601]$

$\text{Range}(S_{0,3}) = [5, 507]$

$\text{Range}(S_{0,4}) = [40, 101]$

$\text{Range}(S_{1,1}) = [2, 201]$

$\text{Range}(S_{1,2}) = [250, 419]$

$\text{Range}(S_{1,3}) = [520, 680]$

$\text{Range}(S_{1,4}) = [708, 1001]$

$\text{Range}(S_{1,5}) = [1040, 1560]$

After
Compaction

$\text{Range}(S_{1,4}) = [708, 1001]$

$\text{Range}(S_{1,5}) = [1040, 1560]$

$\text{Range}(S_{1,6}) = [2, 185]$

$\text{Range}(S_{1,7}) = [199, 240]$

$\text{Range}(S_{1,8}) = [247, 376]$


$\text{Range}(S_{1,9}) = [387, 520]$

$\text{Range}(S_{1,10}) = [543, 680]$


When to trigger leveled compaction?

- Based on size threshold for SSTables
- $Size(L)$ = total size (in MB) of all level- L SSTables
- **Level 0**: Compact when the number of level-0 SSTables reaches a threshold (e.g., 8)
- **Level L , $L \geq 1$** : Compact when $Size(L) > F^L$ MB
 - ▶ $F = 10$ in LevelDB
- Each level stores F times as much data as previous level
 - ▶ $Size(L) \leq F^L$ MB, $L \geq 1$

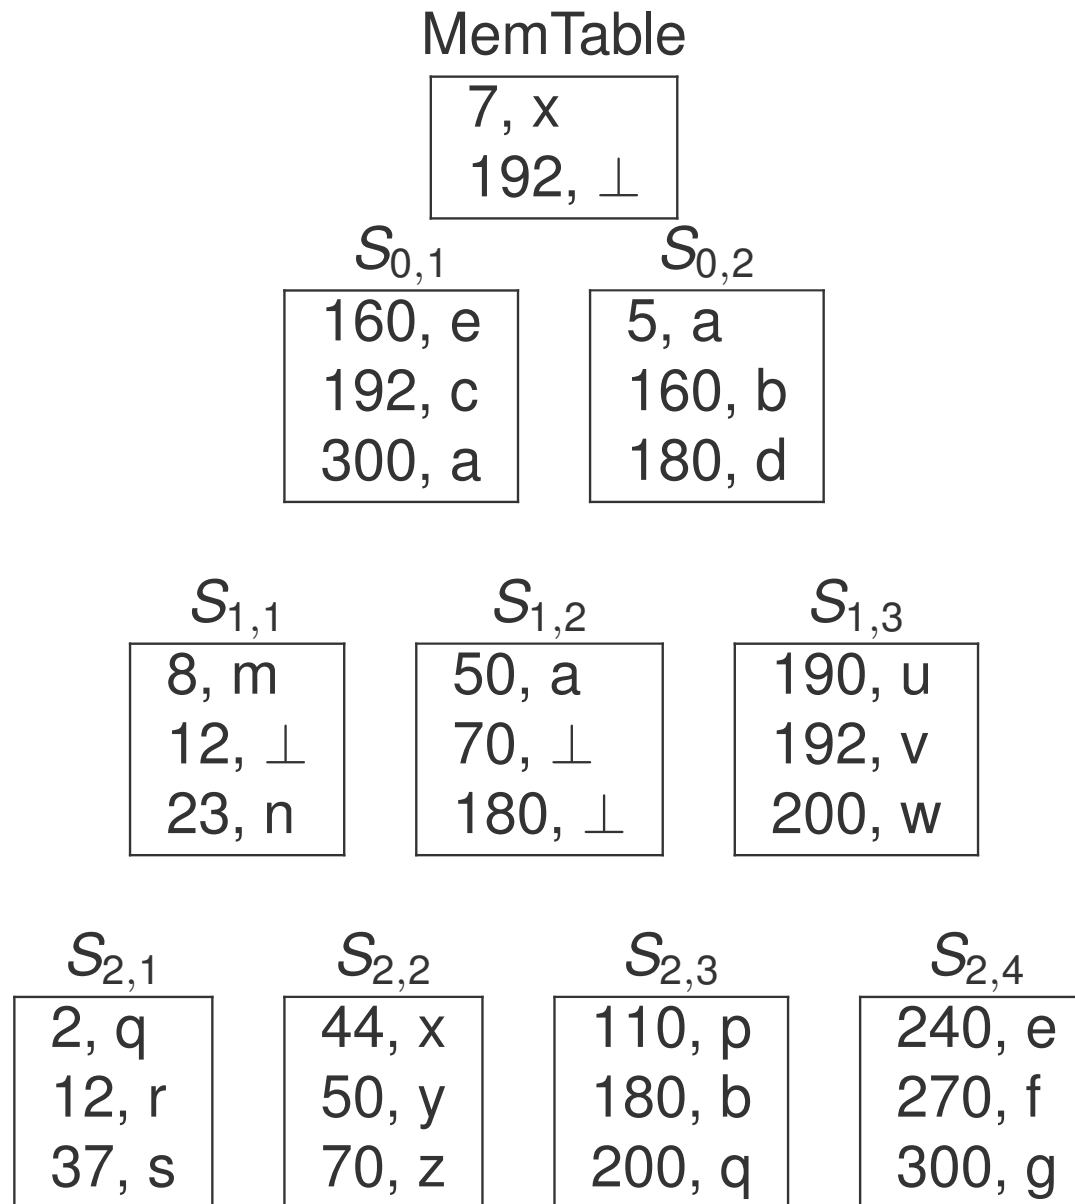
Searching LSM Storage




Search Example 1: search key = 7



Search Example 2: search key = 160



Search Example 3: search key = 200



LCS: Search Algorithm

EqualitySearch (k)

Input: search key k

Output: value of k if found; otherwise, *null*

01. if (k is found in MemTable) then return k 's value
02. let $S_{0,1}, \dots, S_{0,n}$ be the sequence of level-0 SSTables
where $S_{0,i+1}$ is more recent than $S_{0,i}$
03. for $i = n$ downto 1 do
04. if ($k \in \text{Range}(S_{0,i})$) then
05. Search $S_{0,i}$ for k ; if found then return k 's value
06. let m be the maximum number of levels of SSTables
07. for $L = 1$ to m do
08. let $S_{L,1}, S_{L,2}, \dots$ be the sequence of level- L SSTables
09. if there exists i such that $k \in \text{Range}(S_{L,i})$ then
10. Search $S_{L,i}$ for k ; if found then return k 's value
11. return *null*

Optimizing SSTable Search

- Each SSTable is stored as a file consisting of a sequence of data blocks

Block 1	Block 2	Block n-1	Block n
---------	---------	-------	-----------	---------

- How to optimize SSTable search?
 - ▶ Given a SSTable S and search key k , which block in S could contain k ?
 - ▶ Given a block B and search key k , does B contain k ?

Optimization 1: Sparse Index

- Assume each SSTable is 2MB consisting of 512 4KB blocks
- **Problem:** How to quickly locate SSTable block for a given search key?
- **Solution:** Build a sparse index for each SSTable
 - ▶ Sparse index: $(k_1, k_2, \dots, k_{512})$
 - ▶ Each k_i = the first key value in the i^{th} block of SSTable
- **Example:** Consider the following sparse index for a SSTable:

k_1	k_2	k_3	k_4	\dots	k_{512}
5	26	79	204	\dots	8790

To look for key 90 in this SSTable, search the third block

Optimization 2: Bloom Filter

- **Problem:** How to quickly determine whether a search key exists in a SSTable block?
- **Solution:** Build a bloom filter for each block
- **Bloom filter** = Space-efficient randomized data structure for representing a set to support membership queries
 - ▶ B. H. Bloom, *Space/Time Trade-offs in Hash Coding with Allowable Errors*, CACM, 13(7), 422-426, 1970
- Represent a set $S = \{x_1, x_2, \dots, x_n\}$ using a m-bit array, $B[1..m]$
 - ▶ k independent hash functions: h_1, h_2, \dots, h_k
 - ▶ $h_i : S \rightarrow \{1, 2, \dots, m\}$

Optimization 2: Bloom Filter (cont.)

CreateBloomFilter (S, m, h_1, \dots, h_k)

```
01.  Initialize  $B[i] = 0$  for  $i = 1$  to  $m$ 
02.  for  $x \in S$  do
03.      for  $i = 1$  to  $k$  do
04.           $j = h_i(x)$ 
05.          set  $B[j] = 1$ 
06.  return  $B$ 
```

How to use bloom filter to determine if $x \in S$?

- If there exists $i \in [1, k]$ such that $h_i(x) = j$ and $B[j] = 0$, then $x \notin S$
- Otherwise, x could be in S
 - ▶ x is called a **false positive** if x is actually not in S

Optimization 2: Bloom Filter (cont.)

- Build a bloom filter B for $S = \{\text{Curly, Larry, Moe}\}$ with $m=16$ & $k=3$

x	$h_1(x)$	$h_2(x)$	$h_3(x)$
Curly	13	1	4
Larry	5	10	2
Moe	8	2	11

	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1
B	0	0	0	1	0	1	1	0	1	0	0	1	1	0	1	1

x	$h_1(x)$	$h_2(x)$	$h_3(x)$
Alice	4	6	13
Bob	1	10	8

- Based on B , is $\text{Alice} \in S$?
- Based on B , is $\text{Bob} \in S$?

Scatter-gather Queries

- Consider a relation $R(A,B,C)$ that is hash partitioned using attribute A
- Consider two queries on R
 - ▶ Q1: select * from R where $A = 10$ and $B > 20$
 - ▶ Q2: select * from R where $B > 20$
- Q2 is an example of a **scatter-gather query**
 - ▶ Need to access every partition to process query!

Indexing

Customers

cust#	cname	city
1	Alice	Singapore
2	Bob	Jarkata
3	Carol	Bangkok
4	Dave	Jarkata
5	Eve	Singapore
6	Fred	Penang
7	George	Hanoi
8	Hal	Bangkok
9	Ivy	Singapore
10	Joe	Penang
11	Kathy	Singapore
12	Larry	Jarkata

Index on Customers.city

Bangkok	3, 8
Hanoi	7
Jarkata	2, 4, 12
Penang	6, 10
Singapore	1, 5, 9, 11

How to Index Partitioned Data?

Customers₁

cust#	cname	city
3	Carol	Bangkok
6	Fred	Penang
9	Ivy	Singapore
12	Larry	Jarkata

Customers₂

cust#	cname	city
1	Alice	Singapore
4	Dave	Jarkata
7	George	Hanoi
10	Joe	Penang

Customers₃

cust#	cname	city
2	Bob	Jarkata
5	Eve	Singapore
8	Hal	Bangkok
11	Kathy	Singapore

Approach 1: Local Indexing

Customers₁

cust#	cname	city
3	Carol	Bangkok
6	Fred	Penang
9	Ivy	Singapore
12	Larry	Jakarta

Index I_1 on Customers₁.city

Bangkok	3
Jakarta	12
Penang	6
Singapore	9

Customers₂

cust#	cname	city
1	Alice	Singapore
4	Dave	Jakarta
7	George	Hanoi
10	Joe	Penang

Index I_2 on Customers₂.city

Hanoi	7
Jakarta	4
Penang	10
Singapore	1

Customers₃

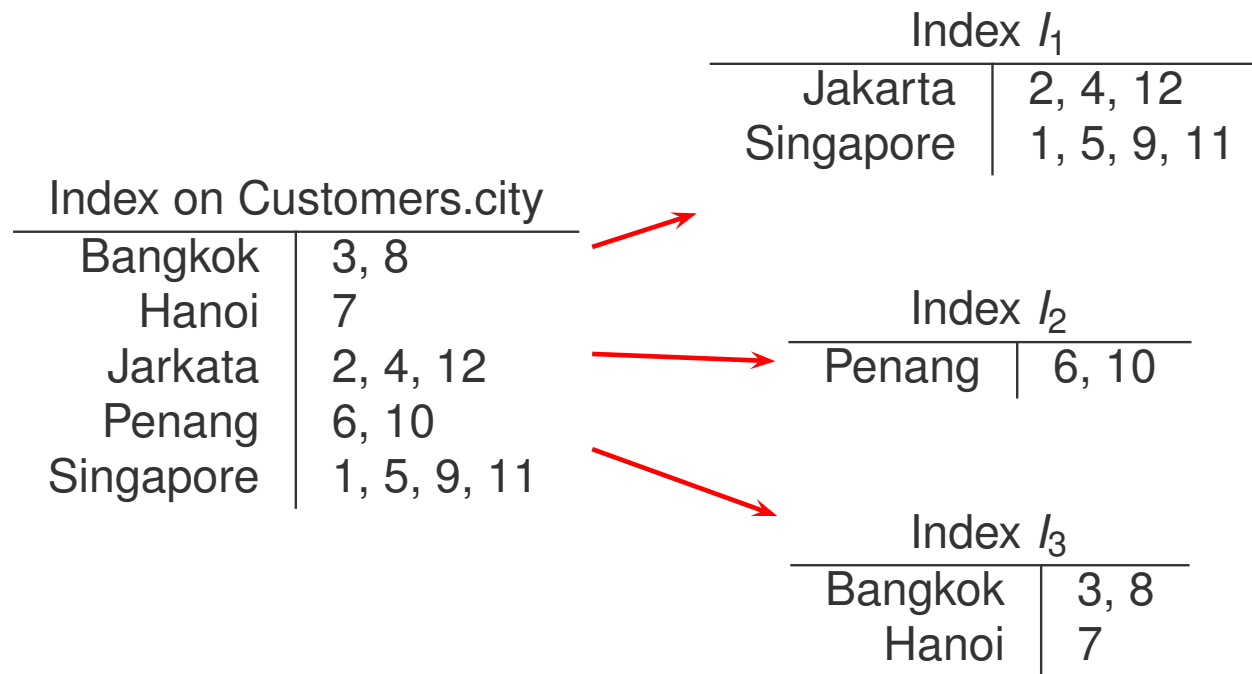
cust#	cname	city
2	Bob	Jakarta
5	Eve	Singapore
8	Hal	Bangkok
11	Kathy	Singapore

Index I_3 on Customers₃.city

Bangkok	8
Jakarta	2
Singapore	5, 11

Approach 2: Global Indexing

city	Hash(city)
Bangkok	3
Hanoi	3
Jakarta	1
Penang	2
Singapore	1



Approach 2: Global Indexing (cont.)

Customers₁

cust#	cname	city
3	Carol	Bangkok
6	Fred	Penang
9	Ivy	Singapore
12	Larry	Jarkata

Index I_1

Jakarta	2, 4, 12
Singapore	1, 5, 9, 11

Customers₂

cust#	cname	city
1	Alice	Singapore
4	Dave	Jarkata
7	George	Hanoi
10	Joe	Penang

Index I_2

Penang	6, 10
--------	-------

Customers₃

cust#	cname	city
2	Bob	Jarkata
5	Eve	Singapore
8	Hal	Bangkok
11	Kathy	Singapore

Index I_3

Bangkok	3, 8
Hanoi	7

Local vs Global Indexing

Partitioned Data

Customers₁

cust#	cname	city
3	Carol	Bangkok
6	Fred	Penang
9	Ivy	Singapore
12	Larry	Jakarta

Local Index

Index I_1 on Customers₁.city

Bangkok	3
Jakarta	12
Penang	6
Singapore	9

Global Index

Index I_1

Jakarta	2, 4, 12
Singapore	1, 5, 9, 11

Customers₂

cust#	cname	city
1	Alice	Singapore
4	Dave	Jakarta
7	George	Hanoi
10	Joe	Penang

Index I_2 on Customers₂.city

Hanoi	7
Jakarta	4
Penang	10
Singapore	1

Index I_2

Penang	6, 10
--------	-------

Customers₃

cust#	cname	city
2	Bob	Jakarta
5	Eve	Singapore
8	Hal	Bangkok
11	Kathy	Singapore

Index I_3 on Customers₃.city

Bangkok	8
Jakarta	2
Singapore	5, 11

Index I_3

Bangkok	3, 8
Hanoi	7

DynamoDB: Data Model

- A **table** is a collection of data
 - ▶ Each table contains zero or more items
- An **item** is a group of attributes that is uniquely identifiable among all of the other items
 - ▶ Each item is composed of one or more **attributes**
- Each item in a table has a unique identifier, or **primary key**
 - ▶ Other than the primary key, each table is schemaless, which means that neither the attributes nor their data types need to be defined beforehand
 - ▶ Each item can have its own distinct attributes
- Two types of primary key
 - ▶ **Simple primary key** = (partition key)
 - ▶ **Composite primary key** = (partition key, sort key)

DynamoDB: Data Model (cont.)

- Each table is partitioned by hashing on the **partition key**
- Items with the same partition key are stored together in sorted order by the **sort key** value

Secondary Indexes

- Base table = table being indexed
- Index key = partition key & (possibly) sort key
- Each index entry contains base table's primary key value & optionally **projected attribute values**
- Two types of secondary indexes
 - ▶ Global Secondary Index (GSI)
 - ▶ Local Secondary Index (LSI)

Global vs Local Secondary index

Global Index

- Index key can be simple or composite
- Partition key could be different from base table's

Local Index

- Index key must be composite
- Partition key must be the same as base table's

Example

- Base table: Customers (cust#, cname, email, city)
 - ▶ Partition key = cust#
- LSI with schema (cust#, city, email)
 - ▶ Partition key = cust#
 - ▶ Sort key = city
 - ▶ Projected attribute = email
- GSI with schema (city, cust#, email)
 - ▶ Partition key = city
 - ▶ Base table's primary key = cust#
 - ▶ Projected attribute = email

References

- S. Ghemawat, J. Dean, LevelDB implementation

<https://github.com/google/leveldb/blob/master/doc/impl.md>

- Cassandra Database Internals: How is data maintained?

https://docs.datastax.com/en/dse/6.8/dse-arch/datastax_enterprise/dbInternals/dbIntHowDataMaintain.html

- Core components of Amazon DynamoDB

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.CoreComponents.html>

- Improving data access with secondary indexes

<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/SecondaryIndexes.html>