

CS4224/CS5424 Lecture 6

Review of Concurrency Control

Transactions

- ▶ A transaction is an abstraction representing a logical unit of work
- ▶ **Example:** `Transfer(x, y, amount)`
 - ▶ transaction to transfer `amount` from account `x` to account `y`

```
BEGIN TRANSACTION;  
SELECT balance INTO :Balx FROM Account WHERE accountId = :x;  
  
SELECT balance INTO :Baly FROM Account WHERE accountId = :y;  
  
If (:Balx < amount) then ABORT;  
  
UPDATE Account SET balance = :Baly + :amount WHERE accountId = :y;  
  
UPDATE Account SET balance = :Balx - :amount WHERE accountId = :x;  
  
COMMIT;
```

Transaction Management

- ▶ Ensures four properties of transactions (Xacts) to maintain data in the face of concurrent access and system failures
 1. **Atomicity**: Either all or none of the actions in Xact happen
 2. **Consistency**: If each Xact is consistent, and the DB starts consistent, the DB ends up consistent
 3. **Isolation**: Execution of one Xact is isolated from other Xacts
 4. **Durability**: If a Xact commits, its effects persist
- ▶ The **concurrency control manager** component ensures isolation
- ▶ The **recovery manager** component ensures atomicity and durability

Transactions

- ▶ A transaction (Xact) T_i can be viewed as a sequence of **actions**:
 - ▶ $R_i(O)$ = T_i reads an object O
 - ▶ $W_i(O)$ = T_i writes an object O
 - ▶ $Commit_i$ = T_i terminates successfully
 - ▶ $Abort_i$ = T_i terminates unsuccessfully
- ▶ Each Xact must end with either a commit or an abort action
- ▶ An **active Xact** is a Xact that is still in progress (i.e., has not yet terminated)

Transactions

- ▶ $R(x), R(y), W(y), W(x)$, Commit
- ▶ $R(x), R(y)$, Abort

```
BEGIN TRANSACTION;  
SELECT balance INTO :Balx FROM Account WHERE accountId = :x;  
  
SELECT balance INTO :Baly FROM Account WHERE accountId = :y;  
  
If (:Balx < amount) then ABORT;  
  
UPDATE Account SET balance = :Baly + :amount WHERE accountId = :y;  
  
UPDATE Account SET balance = :Balx - :amount WHERE accountId = :x;  
  
COMMIT;
```

Transaction Schedules

- ▶ **Schedule** = a list of actions from a set of Xacts, where the order of the actions within each Xact is preserved
- ▶ **Example:** Consider two Xacts T_1 and T_2 :
 - ▶ T_1 : $R_1(A), W_1(A), R_1(B), W_1(B), Commit_1$
 - ▶ T_2 : $R_2(A), W_2(A), R_2(B), W_2(B), Commit_2$
- ▶ Some schedules of T_1 and T_2 :
 - S_1 : $R_1(A), W_1(A), R_1(B), W_1(B), Commit_1, R_2(A), W_2(A), R_2(B), W_2(B), Commit_2$
 - S_2 : $R_2(A), W_2(A), R_2(B), W_2(B), Commit_2, R_1(A), W_1(A), R_1(B), W_1(B), Commit_1$
 - S_3 : $R_1(A), W_1(A), R_2(A), W_2(A), R_1(B), W_1(B), Commit_1, R_2(B), W_2(B), Commit_2$
- ▶ A **serial schedule** is a schedule where the actions of Xacts are not interleaved

Read From & Final Write

- ▶ We say that T_j reads O from T_i in a schedule S if the last write action on O before $R_j(O)$ in S is $W_i(O)$
- ▶ We say that T_j reads from T_i if T_j has read some object from T_i
- ▶ We say that T_i performs the final write on O in a schedule S if the last write action on O in S is $W_i(O)$

- ▶ **Example:** Consider the following schedule:

$R_1(A), W_1(A), R_2(A), W_2(A), R_1(B), W_1(B), \text{Commit}_1, R_2(B), W_2(B), \text{Commit}_2$

- ▶ T_1 reads A from the initial database (or T_1 reads A from T_0)
 - ★ Assume that initial database was created by dummy Xact T_0
- ▶ T_2 reads A from T_1
- ▶ T_1 reads B from T_0
- ▶ T_2 reads B from T_1
- ▶ T_2 performs the final write on A
- ▶ T_2 performs the final write on B

View Serializable Schedules

- ▶ Two schedules S and S' (over the same set of Xacts) are **view equivalent** (denoted by $S \equiv_v S'$) if they satisfy all the following conditions:
 1. If T_i reads A from T_j in S , then T_i must also read A from T_j in S'
 2. For each data object A , the Xact (if any) that performs the final write on A in S must also perform the final write on A in S'
- ▶ **Example:**
 - ▶ S_3 : $R_1(x), R_1(y), W_1(y), R_2(y), R_2(x), W_1(x), W_2(x), W_2(y)$
 - ▶ S_5 : $R_1(x), R_2(y), R_1(y), W_1(y), R_2(x), W_2(x), W_2(y), W_1(x)$
 - ▶ S_6 : $R_1(x), R_2(y), R_1(y), R_2(x), W_1(y), W_2(x), W_1(x), W_2(y)$
 - ▶ $S_3 \not\equiv_v S_5$ as T_2 reads y from T_1 in S_3 but T_2 reads y from T_0 in S_5
 - ▶ Similarly, $S_3 \not\equiv_v S_6$
 - ▶ $S_5 \equiv_v S_6$

View Serializable Schedules (cont.)

- ▶ A schedule S is a **view serializable schedule (VSS)** if S is view equivalent to some serial schedule over the same set of Xacts
- ▶ **Example:**
 - ▶ Consider two Xacts T_1 and T_2 :
 - ★ T_1 : $R_1(x), R_1(y), W_1(y), W_1(x)$
 - ★ T_2 : $R_2(y), R_2(x), W_2(x), W_2(y)$
 - ▶ Serial schedules over $\{T_1, T_2\}$:
 - ★ S_1 : $R_1(x), R_1(y), W_1(y), W_1(x), R_2(y), R_2(x), W_2(x), W_2(y)$
 - ★ S_2 : $R_2(y), R_2(x), W_2(x), W_2(y), R_1(x), R_1(y), W_1(y), W_1(x)$
 - ▶ S_3 is not view serializable; S_4 is view serializable
 - ★ S_3 : $R_1(x), R_1(y), W_1(y), R_2(y), R_2(x), W_1(x), W_2(x), W_2(y)$
 - ★ S_4 : $R_1(x), R_1(y), W_1(y), R_2(y), W_1(x), R_2(x), W_2(x), W_2(y)$

Conflicting Actions

- ▶ Two actions on the same object **conflict** if
 1. at least one of them is a write action, and
 2. the actions are from different Xacts
- ▶ **Examples:**
 - ▶ $R_1(x)$ and $R_2(x)$ do not conflict
 - ▶ $R_1(x)$ and $W_1(x)$ do not conflict
 - ▶ $W_1(x)$ and $R_2(y)$ do not conflict
 - ▶ $W_1(x)$ and $R_2(x)$ conflict
 - ▶ $W_1(x)$ and $W_2(x)$ conflict

Conflict Serializable Schedules


- ▶ Two schedules S & S' (over the same set of Xacts) are said to be **conflict equivalent** (denoted by $S \equiv_c S'$) if they order every pair of conflicting actions of two committed Xacts in the same way
- ▶ **Example:**
 - ▶ S_3 : $R_1(x)$, $R_1(y)$, $W_1(y)$, $R_2(y)$, $R_2(x)$, $W_1(x)$, $W_2(x)$, $W_2(y)$
 - ▶ S_5 : $R_1(x)$, $R_2(y)$, $R_1(y)$, $W_1(y)$, $R_2(x)$, $W_2(x)$, $W_2(y)$, $W_1(x)$
 - ▶ S_6 : $R_1(x)$, $R_2(y)$, $R_1(y)$, $R_2(x)$, $W_1(y)$, $W_2(x)$, $W_1(x)$, $W_2(y)$
 - ▶ $S_3 \not\equiv_c S_5$ as $W_1(y)$ precedes $R_2(y)$ in S_3 but not in S_5
 - ▶ Similarly, $S_3 \not\equiv_c S_6$
 - ▶ $S_5 \equiv_c S_6$

Conflict Serializable Schedules (cont.)

- ▶ A schedule is a **conflict serializable schedule (CSS)** if it is conflict equivalent to a serial schedule over the same set of Xacts
- ▶ **Example:**
 - ▶ Consider two Xacts T_1 and T_2 :
 - ★ T_1 : $R_1(x)$, $R_1(y)$, $W_1(y)$, $W_1(x)$
 - ★ T_2 : $R_2(y)$, $R_2(x)$, $W_2(x)$, $W_2(y)$
 - ▶ Serial schedules over $\{T_1, T_2\}$:
 - ★ S_1 : $R_1(x)$, $R_1(y)$, $W_1(y)$, $W_1(x)$, $R_2(y)$, $R_2(x)$, $W_2(x)$, $W_2(y)$
 - ★ S_2 : $R_2(y)$, $R_2(x)$, $W_2(x)$, $W_2(y)$, $R_1(x)$, $R_1(y)$, $W_1(y)$, $W_1(x)$
 - ▶ S_3 is not conflict serializable; S_4 is conflict serializable
 - ★ S_3 : $R_1(x)$, $R_1(y)$, $W_1(y)$, $R_2(y)$, $R_2(x)$, $W_1(x)$, $W_2(x)$, $W_2(y)$
 - ★ S_4 : $R_1(x)$, $R_1(y)$, $W_1(y)$, $R_2(y)$, $W_1(x)$, $R_2(x)$, $W_2(x)$, $W_2(y)$

Testing for Conflict Serializability

- ▶ A **conflict serializability graph** for a schedule S (denoted by $CSG(S)$) is a directed graph $G = (V, E)$ such that
 - ▶ V contains a node for each committed Xact in S
 - ▶ E contains (T_i, T_j) if an action in T_i precedes and conflicts with one of T_j 's actions
- ▶ **Example:** Conflict serializability graph for schedule $R_1(A), W_2(A), Commit_2, W_1(A), Commit_1, W_3(A), Commit_3$



- ▶ **Theorem 1:** A schedule is conflict serializable iff its conflict serializability graph is acyclic
- ▶ **Theorem 2:** A schedule that is conflict serializable is also view serializable

Blind Writes

- ▶ A write on object O by T_i is call a **blind write** if T_i did not read O prior to the write
 - ▶ Consider the schedule: $R_1(x), W_2(y), W_1(x)$
 - ▶ $W_2(y)$ is a blind write
 - ▶ $W_1(x)$ is a non-blind write
- ▶ **Theorem 3:** If S is view serializable and S has no blind writes, then S is also conflict serializable

Recoverable Schedules

T_1	T_2
$W_1(x)$	$R_2(x)$ $W_2(y)$ <i>Commit₂</i>

Non-Recoverable
Schedule

- ▶ A schedule S is said to be a **recoverable schedule** if for every Xact T that commits in S , T must commit after T' if T reads from T'

Lock-Based Concurrency Control

- ▶ Each Xact needs to request for an appropriate **lock** on an object before the Xact can access the object
- ▶ **Locking modes**
 - ▶ **Shared (S) locks** for reading objects
 - ▶ **Exclusive (X) locks** for writing objects
- ▶ **Lock compatibility:**

Lock Requested	Lock Held		
	-	S	X
S	✓	✓	×
X	✓	×	×

✓: Compatible
Lock request is granted

×: Incompatible
Lock request is blocked

Lock-Based Concurrency Control (cont.)

1. To **read an object** O , a Xact must request for a shared/exclusive lock on O
2. To **update an object** O , a Xact must request for an exclusive lock on O
3. A **lock request is granted** on O if the requesting lock mode is compatible with the lock modes of existing locks on O
4. If T 's **lock request is not granted** on O , T becomes **blocked**: its execution is suspended & T is added to O 's **request queue**
5. When a **lock is released** on O , the lock manager checks the request of the first Xact T in the request queue for O . If T 's request can be granted, T acquires its lock on O and resumes execution after its removal from the queue
6. When a Xact **commits/aborts**, all its locks are released & T is removed from any request queue it's in

Two Phase Locking (2PL) Protocol

► 2PL Protocol:

1. To read an object O, a Xact must hold a S-lock or X-lock on O
2. To write to an object O, a Xact must hold a X-lock on O
3. Once a Xact releases a lock, the Xact can't request any more locks

► Xacts using 2PL can be characterized into two phases:

- Growing phase: before releasing 1st lock
- Shrinking phase: after releasing 1st lock

► **Theorem 4:** 2PL schedules are conflict serializable

Strict Two Phase Locking (strict 2PL) Protocol

► 2PL Protocol:

1. To read an object O, a Xact must hold a S-lock or X-lock on O
2. To write to an object O, a Xact must hold a X-lock on O
3. Once a Xact releases a lock, the Xact can't request any more locks

► **Theorem 4:** 2PL schedules are conflict serializable

► Strict 2PL Protocol:

1. To read an object O, a Xact must hold a S-lock or X-lock on O
2. To write to an object O, a Xact must hold a X-lock on O
3. A Xact must hold on to locks until Xact commits or aborts

► **Theorem 5:** Strict 2PL schedules are strict & conflict serializable

Lock Management

- ▶ Handling deadlocks
- ▶ Lock conversion


Deadlocks

- ▶ **Deadlock**: cycle of Xacts waiting for locks to be released by each other
- ▶ **Example**:
 - T_1 requests X-lock on A and is granted;
 - T_2 requests X-lock on B and is granted;
 - T_1 requests X-lock on B and is blocked;
 - T_2 requests X-lock on A and is blocked;
- ▶ Dealing with deadlocks:
 - ▶ deadlock detection
 - ▶ deadlock prevention

How to Detect Deadlocks?

► Waits-for graph (WFG)

- Nodes represent active Xacts
- Add an edge $T_i \rightarrow T_j$ if T_i is waiting for T_j to release a lock



► Lock manager

- adds an edge when it queues a lock request
- updates edges when it grants a lock request

► Deadlock is detected if WFG has a cycle

► Breaks a deadlock by aborting a Xact in cycle

► Alternative to WFG: timeout mechanism

How to Prevent Deadlocks?

- ▶ Assume older Xacts have higher priority than younger Xacts
 - ▶ Each Xact is assigned a timestamp when it starts
 - ▶ An older Xact has a smaller timestamp
- ▶ Suppose T_i requests for a lock that conflicts with a lock held by T_j
- ▶ Two possible deadlock prevention policies:

Prevention Policy	T_i has higher priority	T_i has lower priority
Wait-die	T_i waits for T_j	T_i aborts
Wound-wait	T_j aborts	T_i waits for T_j

- ▶ **Wait-die policy**
 - ▶ non-preemptive: only a Xact requesting for a lock can get aborted
 - ▶ a younger Xact may get repeatedly aborted
 - ▶ a Xact that has all the locks it needs is never aborted
- ▶ **Wound-wait policy**
 - ▶ preemptive
- ▶ To avoid starvation, a restarted Xact must use its original timestamp!

Lock Conversion

- ▶ Consider two Xacts:

$T_1: R_1(A), R_1(B), W_1(A)$

$T_2: R_2(A), R_2(B)$

- ▶ Since T_1 needs to both read & update A , T_1 acquires an exclusive lock to read A
- ▶ All possible schedules are serial

$S_1: R_1(A), R_1(B), W_1(A), R_2(A), R_2(B)$

$S_2: R_2(A), R_2(B), R_1(A), R_1(B), W_1(A)$

- ▶ Increase concurrency by allowing **lock conversions**
- ▶ Two types of lock conversions
 - ▶ T_i **upgrades** its S-lock on object A to X-lock
 - ▶ T_i **downgrades** its X-lock on object A to S-lock

Lock Conversion (cont.)

- ▶ Interleaved executions become possible with lock upgrading:

S_3 : $R_1(A)$, $R_2(A)$, $R_2(B)$, $R_1(B)$, $W_1(A)$

S_4 : $R_1(A)$, $R_1(B)$, $R_2(A)$, $R_2(B)$, $W_1(A)$

S_5 : $R_1(A)$, $R_2(A)$, $R_1(B)$, $R_2(B)$, $W_1(A)$

- ▶ **Lock upgrade** on object A

- ▶ Upgrade request is blocked if another Xact is holding a shared lock on A
- ▶ Upgrade request is allowed if T_i has not released any lock

- ▶ **Lock downgrade** on object A

- ▶ Downgrade request is allowed if
 1. T_i has not modified A , and
 2. T_i has not released any lock

Multiversion Concurrency Control (MVCC)

T_1	T_2
$R_1(x)$	
$W_1(x)$	
	$R_2(x)$
	$W_2(y)$
$R_1(y)$	
$W_1(z)$	

Schedule S

T_1	T_2
$R_1(x)$	
$W_1(x)$	
$R_1(y)$	
	$R_2(x)$
	$W_2(y)$
$W_1(z)$	

Schedule S'

Multiversion Concurrency Control (MVCC)

- ▶ **Key idea:** maintain multiple versions of each object
 - ▶ $W_i(O)$ creates a new version of object O
 - ▶ $R_i(O)$ reads an appropriate version of O
- ▶ Advantages:
 - ▶ Read-only Xacts are not blocked by update Xacts
 - ▶ Update Xacts are not blocked by read-only Xacts
 - ▶ Read-only Xacts are never aborted
- ▶ Notation:
 - ▶ $W_i(x)$ creates a new version of x denoted by x_i
 - ▶ For each object x , its initial version is denoted by x_0

MVCC: Example 1

T_1	T_2	Comments
		$x_0 = 10$
	$R_2(x)$	10
$W_1(x)$?

- ▶ In **2PL**, $W_1(x)$ will be blocked
- ▶ In **MVCC**,
 - ▶ T_1 creates a new version of x
 - ▶ Update Xacts are not blocked by read-only Xacts

MVCC: Example 2

T_1	T_2	Comments
$W_1(y)$		$x_0 = 10, y_0 = 20$
	$R_2(x)$	10
	$R_2(y)$	$y_1 = 100$?

- ▶ In **2PL**, $R_2(y)$ will be blocked
- ▶ In **MVCC**,
 - ▶ $W_1(y)$ creates a new version of y (with a value of 100)
 - ▶ $R_2(y)$ is not blocked
 - ▶ $R_2(y)$ returns 20 (the value of the version before $W_1(y)$)
 - ▶ Read-only Xacts are never blocked/aborted

Multiversion Schedules

- ▶ If there are multiple versions of an object x , a read action on x could return any version
- ▶ Thus, an interleaved execution could correspond to different multiversion schedules depending on the MVCC protocol
- ▶ **Example:** $R_1(x)$, $W_1(x)$, $R_2(x)$, $W_2(y)$, $R_1(y)$, $W_1(z)$

S_1 : $R_1(x_0)$, $W_1(x_1)$, $R_2(x_0)$, $W_2(y_2)$, $R_1(y_0)$, $W_1(z_1)$
 S_2 : $R_1(x_0)$, $W_1(x_1)$, $R_2(x_0)$, $W_2(y_2)$, $R_1(y_2)$, $W_1(z_1)$
 S_3 : $R_1(x_0)$, $W_1(x_1)$, $R_2(x_1)$, $W_2(y_2)$, $R_1(y_0)$, $W_1(z_1)$
 S_4 : $R_1(x_0)$, $W_1(x_1)$, $R_2(x_1)$, $W_2(y_2)$, $R_1(y_2)$, $W_1(z_1)$

- ▶ $R_1(x)$ returns x_0
- ▶ $R_2(x)$ could return x_0 or x_1
- ▶ $R_1(y)$ could return y_0 or y_2

Multiversion View Equivalence

- ▶ Two schedules, S and S' , over the same set of transactions, are defined to be **multiversion view equivalent** ($S \equiv_{mv} S'$) if they have the same set of **read-from relationships**
 - ▶ i.e. $R_i(x_j)$ occurs in S iff $R_i(x_j)$ occurs in S'
- ▶ **Example:**
 - S_1 : $R_3(x_0)$, $W_3(x_3)$, $Commit_3$, $W_1(x_1)$, $Commit_1$, $R_2(x_1)$, $W_2(y_2)$, $Commit_2$
 - S_2 : $R_3(x_0)$, $W_3(x_3)$, $Commit_3$, $W_1(x_1)$, $R_2(x_3)$, $Commit_1$, $W_2(y_2)$, $Commit_2$
 - S_3 : $W_1(x_1)$, $Commit_1$, $R_2(x_1)$, $R_3(x_0)$, $W_2(y_2)$, $W_3(x_3)$, $Commit_3$, $Commit_2$
 - ▶ $S_1 \not\equiv_{mv} S_2$ because $R_2(x_1) \in S_1$ and $R_2(x_3) \in S_2$
 - ▶ $S_1 \equiv_{mv} S_3$

Monoversion Schedules

- ▶ A multiversion schedule S is called a **monoversion schedule** if each read action in S returns the most recently created object version
- ▶ **Example:** $R_1(x), W_1(x), R_2(x), W_2(y), R_1(y), W_1(z)$
 - $S_1:$ $R_1(x_0), W_1(x_1), R_2(x_0), W_2(y_2), R_1(y_0), W_1(z_1)$
 - $S_2:$ $R_1(x_0), W_1(x_1), R_2(x_0), W_2(y_2), R_1(y_2), W_1(z_1)$
 - $S_3:$ $R_1(x_0), W_1(x_1), R_2(x_1), W_2(y_2), R_1(y_0), W_1(z_1)$
 - $S_4:$ $R_1(x_0), W_1(x_1), R_2(x_1), W_2(y_2), R_1(y_2), W_1(z_1)$
 - ▶ S_4 is a monoversion schedule
 - ▶ $S_1, S_2,$ and S_3 are not monoversion schedules

Serial Monoversion Schedules

- ▶ A monoversion schedule is defined to be a **serial monoversion schedule** if it is also a serial schedule

- ▶ **Example:**

$S_1:$ $R_1(x_0)$, $W_1(x_1)$, $R_2(x_1)$, $W_2(y_2)$, $R_1(y_2)$, $W_1(z_1)$

$S_2:$ $R_1(x_0)$, $W_1(x_1)$, $R_1(y_0)$, $W_1(z_1)$, $R_2(x_1)$, $W_2(y_2)$

- ▶ S_1 is a non-serial monoversion schedule
- ▶ S_2 is a serial monoversion schedule

Multiversion View Serializability

A multiversion schedule S is defined to be **multiversion view serializable schedule (MVSS)** if there exists a **serial monoversion schedule** (over the same set of Xacts) that is multiversion view equivalent to S

MVSS: Example 1

- ▶ Consider schedule S :

$W_1(x_1), \text{Commit}_1, R_2(x_1), R_3(x_0), W_2(y_2), W_3(x_3), \text{Commit}_3, \text{Commit}_2$

- ▶ S is multiversion view serializable as $S \equiv_{mv} (T_3, T_1, T_2)$:

$R_3(x_0), W_3(x_3), \text{Commit}_3, W_1(x_1), \text{Commit}_1, R_2(x_1), W_2(y_2), \text{Commit}_2$

MVSS: Example 2

- Consider the following schedule S :

T_1	T_2
$R_1(x_0)$	
	$R_2(x_0)$
$R_1(y_0)$	
	$R_2(y_0)$
$W_1(x_1)$	
$Commit_1$	
	$W_2(y_2)$
	$Commit_2$

- S is not multiversion view serializable
- S is not multiversion view equivalent to any serial monoversion schedule
 - $R_1(x_0), R_1(y_0), W_1(x_1), C_1, R_2(x_1), R_2(y_0), W_2(y_2), C_2$
 - $R_2(x_0), R_2(y_0), W_2(y_2), C_2, R_1(x_0), R_1(y_2), W_1(x_1), C_1$

MVSS: Example 3

- Consider the following schedule S :

T_1	T_2	T_3
$R_1(y_0)$	$R_2(x_0)$	
$W_1(y_1)$	$R_2(y_0)$	
$Commit_1$	$W_2(x_2)$	
		$R_3(x_0)$
		$R_3(y_1)$
		$Commit_3$
	$Commit_2$	

- S is not multiversion view serializable
 - Suppose S' is a serial monoversion schedule where $S' \equiv_{mv} S$
 - T_3 must precede T_2 in S' due to $R_3(x_0)$ & $W_2(x_2)$
 - T_2 must precede T_1 in S' due to $R_2(y_0)$ & $W_1(y_1)$
 - T_1 must precede T_3 in S' due to $W_1(y_1)$ & $R_3(y_1)$

Multiversion View Serializability

- ▶ **Theorem 6:** A view serializable schedule (VSS) is also a multiversion view serializable schedule (MVSS)
- ▶ A MVSS is not necessarily VSS
- ▶ **Example:**

$$\begin{array}{ll} T_1: & R_1(x_0), \quad R_1(y_0), \textit{Commit}_1 \\ T_2: & W_2(x_2), W_2(y_2), \textit{Commit}_2, \end{array}$$

- ▶ The above schedule is multiversion view equivalent to the serial monoversion schedule (T_1, T_2)
- ▶ However, the schedule is not a valid monoversion schedule (due to $W_2(y_2)$ & $R_1(y_0)$) and is therefore not VSS

MVCC Protocols

- ▶ Multiversion two-phase locking
- ▶ Multiversion timestamp ordering
- ▶ Snapshot isolation

Snapshot Isolation (SI)

- ▶ Widely used (e.g., Oracle, PostgreSQL, SQL Server, Sybase IQ)
- ▶ Each Xact T sees a snapshot of DB that consists of updates by Xacts that committed before T starts
- ▶ Each Xact T is associated with two timestamps:
 - ▶ $\text{start}(T)$: the time that T starts
 - ▶ $\text{commit}(T)$: the time that T commits

Concurrent Transactions

- ▶ Two Xacts T and T' are defined to be **concurrent** if they overlap
 - ▶ i.e., $[start(T), commit(T)] \cap [start(T'), commit(T')] \neq \emptyset$
- ▶ **Example:**

Timestamp	T_1	T_2	T_3
1	$R_1(B)$		
2		$R_2(A)$	
3	$W_1(B)$		
4	$Commit_1$		
5		$R_2(B)$	
6		$W_2(A)$	
7			$R_3(A)$
8			$R_3(B)$
9			$Commit_3$
10		$Commit_2$	

Snapshot Isolation (SI)

- ▶ $W_i(O)$ creates a version of O denoted by O_i
- ▶ O_i is a **more recent (or newer) version** compared to O_j if $commit(T_i) > commit(T_j)$
- ▶ $R_i(O)$ reads either its own update (if $W_i(O)$ precedes $R_i(O)$) or the latest version of O that is created by a Xact that committed before T_i started; i.e., If $R_i(O)$ returns O_j , then
 1. Either $j = i$ if $W_i(O)$ precedes $R_i(O)$;
 2. Or
 - 2.1 $commit(T_j) < start(T_i)$, and
 - 2.2 For every Xact T_k , $k \neq j$, that has created a version O_k of O , if $commit(T_k) < start(T_i)$, then $commit(T_k) < commit(T_j)$

Example

T_1	T_2	T_3	Comments
$R_1(x)$			x_0
$W_1(x)$			x_1
$R_1(y)$			y_0
	$R_2(x)$		x_0
$W_1(y)$			y_1
$Commit_1$			
	$R_2(y)$		y_0
	$W_2(x)$		x_2
		$R_3(x)$	x_1
		$R_3(y)$	y_1
		$W_3(y)$	y_3
		$R_3(y)$	y_3
		$Commit_3$	

Snapshot Isolation

- ▶ **Concurrent Update Property:** If multiple concurrent Xacts updated the same object, only one of Xacts is allowed to commit
- ▶ If not, the schedule may not be serializable
- ▶ **Example:** Consider the following schedule S

T_1 :	$R_1(x_0)$	$W_1(x_1)$	$Commit_1$
T_2 :	$R_2(x_0)$	$W_2(x_2)$	$Commit_2$

S is not serializable!

- ▶ Two approaches to enforce the concurrent update property:
 - ▶ First Committer Wins (FCW) Rule
 - ▶ First Updater Wins (FUW) Rule

First Committer Wins (FCW) Rule

- ▶ Before committing a Xact T , the system checks if there exists a committed concurrent Xact T' that has updated some object that T has also updated
- ▶ If T' exists, then T aborts
- ▶ Otherwise, T commits

- ▶ **Example 1:**

T_1 :	$R_1(x)$		$W_1(x)$		$Abort_1$
T_2 :		$R_2(x)$		$W_2(x)$	$Commit_2$

- ▶ **Example 2:**

T_1 :	$R_1(x)$		$W_1(x)$		$Commit_1$
T_2 :		$R_2(x)$		$W_2(x)$	$Abort_2$

First Updater Wins (FUW) Rule

- ▶ Whenever a Xact T needs to update an object O , T requests for a X-lock on O
- ▶ If the X-lock is not held by any concurrent Xact, then
 - ▶ T is granted the X-lock on O
 - ▶ If O has been updated by any concurrent Xact, then T aborts
 - ▶ Otherwise, T proceeds with its execution
- ▶ Otherwise, if the X-lock is being held by some concurrent Xact T' , then T waits until T' aborts or commits
 - ▶ If T' aborts, then
 - ★ Assume that T is granted the X-lock on O
 - ★ If O has been updated by any concurrent Xact, then T aborts
 - ★ Otherwise, T proceeds with its execution
 - ▶ If T' commits, then T is aborted
- ▶ When a Xact commits/aborts, it releases its X-lock(s)

Write Skew Anomaly

T_1	T_2
$R_1(x_0)$	$R_2(x_0)$
$R_1(y_0)$	$R_2(y_0)$
$W_1(x_1)$	
$Commit_1$	$W_2(y_2)$
	$Commit_2$

The above is a SI schedule that is not a MVSS

Read-Only Transaction Anomaly

T_1	T_2	T_3
$R_1(y_0)$	$R_2(x_0)$	
$W_1(y_1)$		
$Commit_1$	$R_2(y_0)$	
	$W_2(x_2)$	
		$R_3(x_0)$
		$R_3(y_1)$
		$Commit_3$
	$Commit_2$	

The above is a SI schedule that is not a MVSS

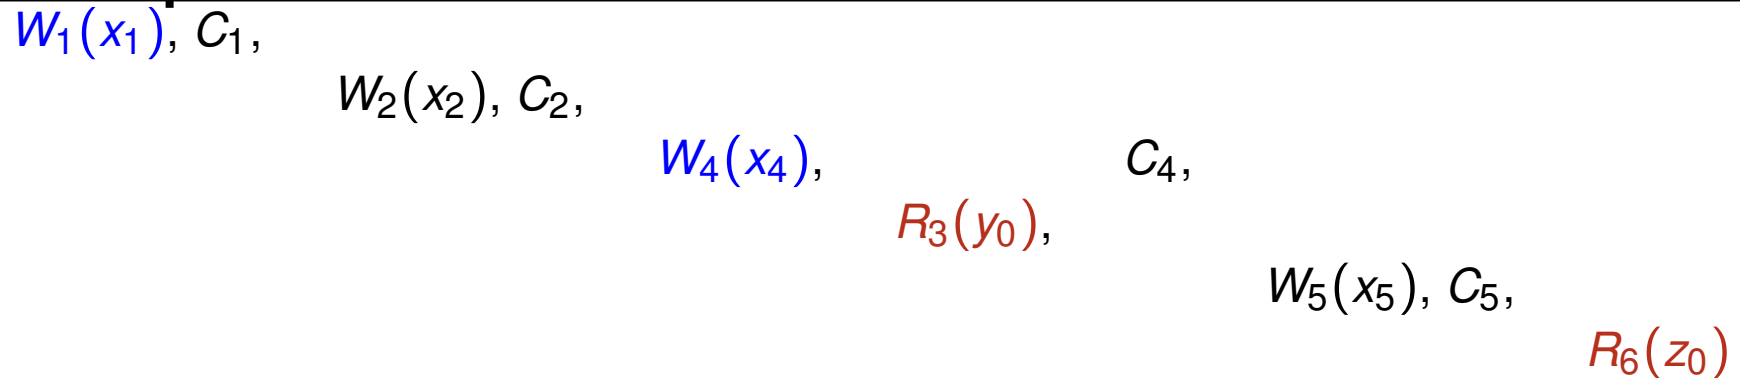
Serializable Snapshot Isolation (SSI) Protocol

- ▶ This is a stronger protocol that guarantees serializable SI schedules
- ▶ A schedule S is defined to be a **serializable snapshot isolation (SSI) schedule** if S is produced by the snapshot isolation protocol (i.e., S is a SI schedule) and S is MVSS

Garbage Collection

- ▶ A version O_i of object O may be deleted if there exists a newer version O_j (i.e., $\text{commit}(T_i) < \text{commit}(T_j)$) such that for every active Xact T_k that started after the commit of T_i (i.e., $\text{commit}(T_i) < \text{start}(T_k)$), we have $\text{commit}(T_j) < \text{start}(T_k)$

- ▶ **Example:**



- ▶ Active transactions: T_3 & T_6
- ▶ Versions that can be deleted: x_1 & x_4