


# CS4224/CS5424 Lecture 9

## Raft Consensus Algorithm



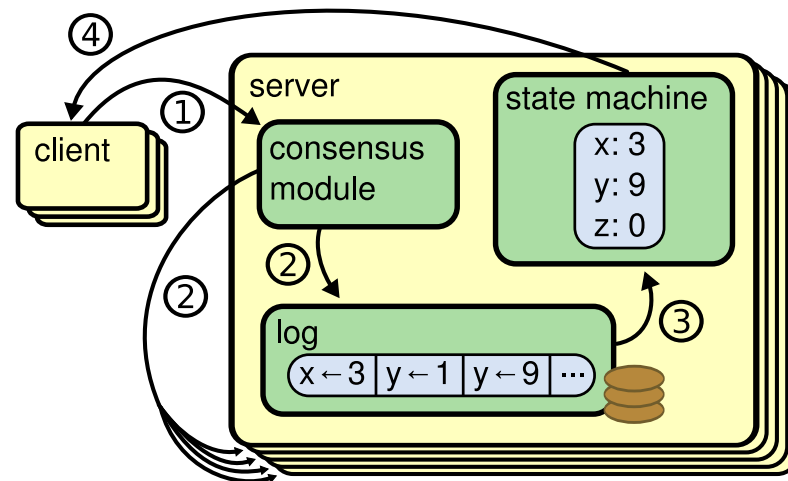
(Ongaro & Ousterhout, 2014)

# Consensus Algorithms

- **Consensus Problem:** How to get multiple servers to agree on the same state
- **Consensus Algorithms:**
  - ▶ Viewstamped Replication, 1988
  - ▶ Paxos, 1990
  - ▶ Zab (Zookeeper Atomic Broadcast), 2011
  - ▶ Raft, 2014

# Replicated State Machines (RSM)

- Consensus algorithm used to implement RSM to provide **fault-tolerant** distributed services
- Service is available as long as a majority of servers are operational and can communicate with each other and clients
- **Consensus algorithm** ensures each RSM receives the same sequence of inputs




(Ongaro & Ousterhout, 2014)

# Raft Concepts

- Server states: follower, candidate, leader
- Term
- Log
  - ▶ Log entry: (index, term, command)
  - ▶ Log comparison - determine which log is more up-to-date (a.k.a. more complete)
  - ▶ Committed log entry
- Remote Procedure Calls (RPCs):
  - ▶ RequestVote
  - ▶ AppendEntries
- Leader election

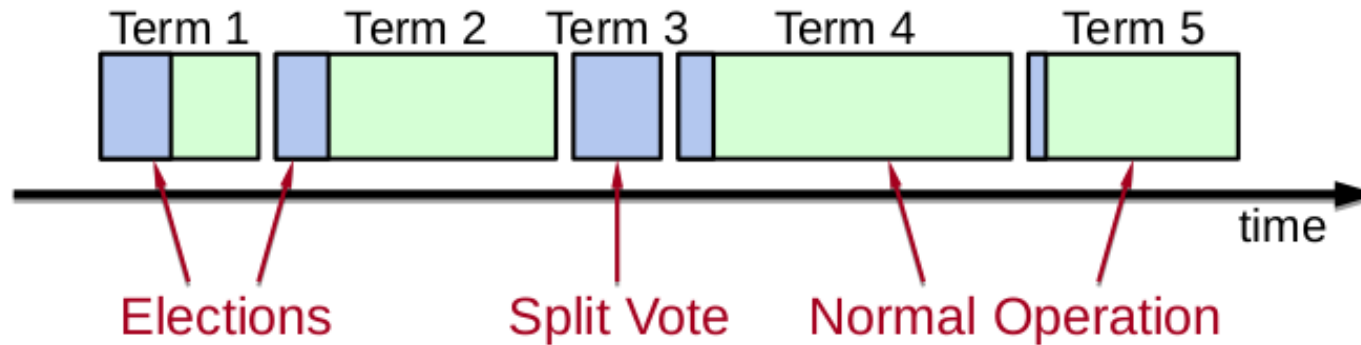
# Server States



(Ongaro & Ousterhout, 2014)

- **Follower** - Passive but expects regular heartbeats from leader
- **Candidate** - Issues **RequestVote RPCs** to get elected as leader
- **Leader** - handles client interactions & issues **AppendEntries RPCs**
  - ▶ Replicate its log
  - ▶ Sends heartbeats to maintain leadership


# Terms



(Ongaro & Ousterhout, 2014)

- Each term starts with an election
- At most one leader elected in each term
- Each server maintains **current term** value
- RPCs/replies include sender's current term
- Server updates its current term number if it receives a message with larger term number

# Logs



(Ongaro & Ousterhout, 2014)

# Remote Procedure Calls (RPCs)

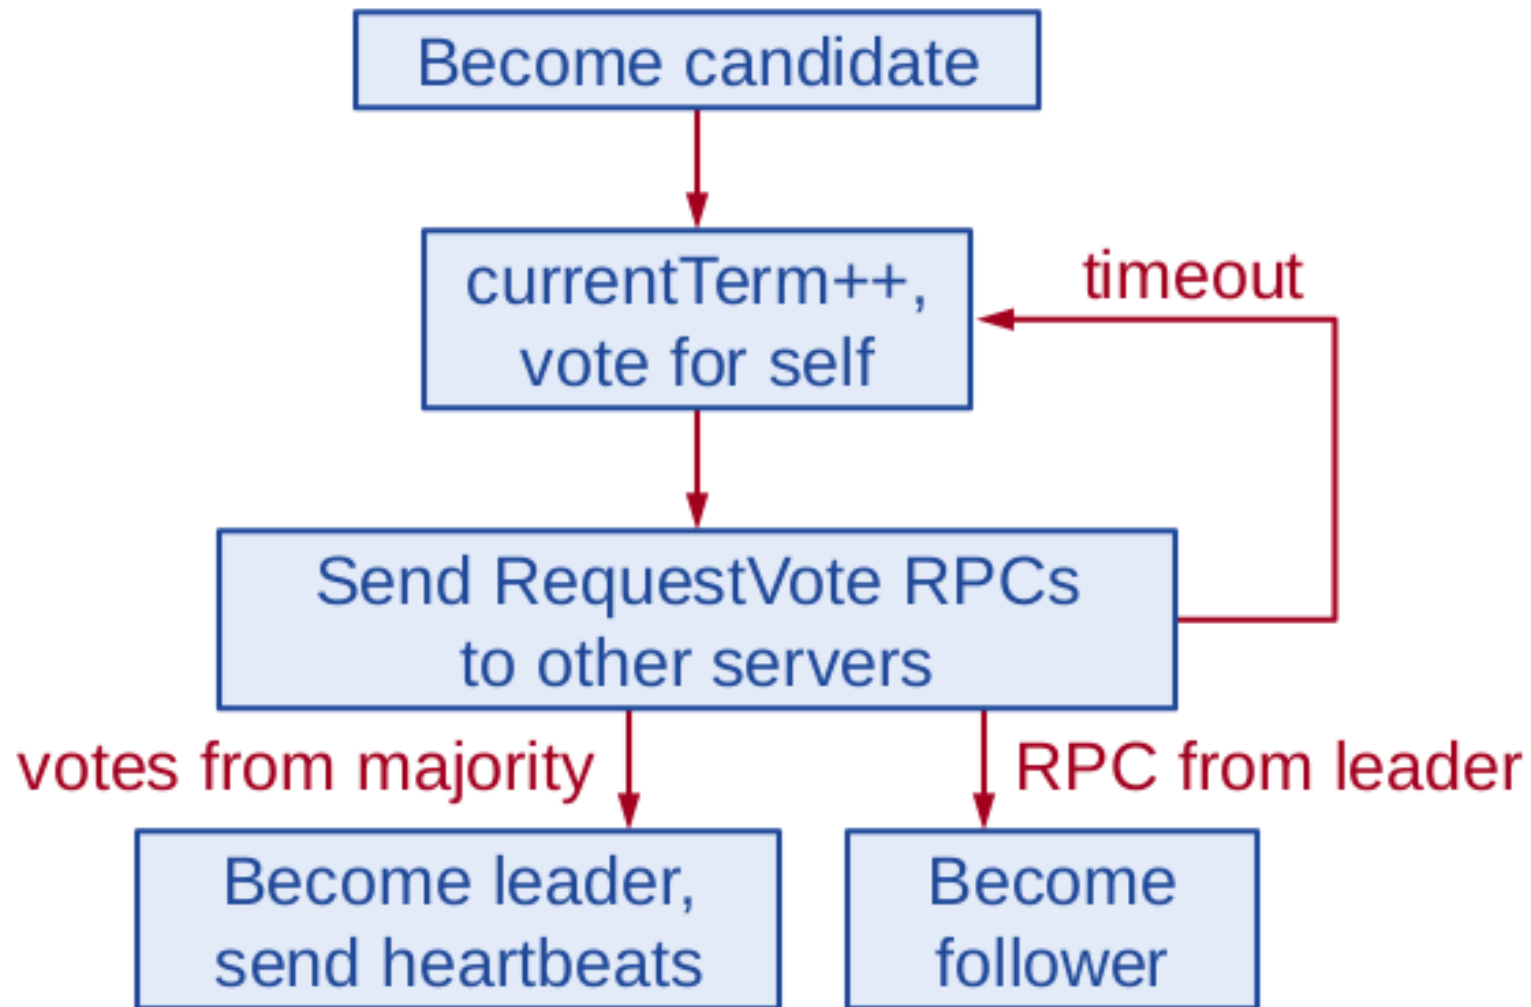
- **RequestVote RPC**
  - ▶ Send by candidate to request for votes to be elected as leader
- **AppendEntries RPC**
  - ▶ Send by leader to replicate its log or as a heartbeat message
- A RPC is resent to server R if leader didn't receive R's response when leader timer timeouts



# Timers

- **Election timer** - follower becomes a candidate or candidate restarts a new election if it didn't receive any RPC
- **Leader timer** - leader resends RPC to follower F if leader didn't receive F's response
- **Client timer** - client resends command if it didn't receive leader's response to command

# Leader Election



(Ongaro & Ousterhout, 2014)

# Election Properties

- **Election Safety Property:** at most one leader can be elected in any term
  - ▶ Each server gives out only one vote each term
  - ▶ A candidate becomes elected as a leader if it receives a majority of votes
- **Election Liveness Property:** some leader must eventually be elected
  - ▶ Each server chooses an election timeout duration randomly from  $[T, 2T]$
  - ▶ Works well if  $T \gg$  broadcast time
    - ★ Broadcast time = average time for server to send RPCs and receive their responses
  - ▶ One server usually wins election before other election timers timeout

# Persistent State on All Servers

- **currentTerm** - latest term that server has seen
- **votedFor** - candidate that received vote in current term (null if none)
- **log[]** - log entries of the form (index,term,command). First index is 1.

# RequestVote RPC

- RequestVote RPC Arguments:
  - ▶ **candidateId** = identifier of candidate
  - ▶ **term** = candidate's term
  - ▶ **lastLogIndex** = index of candidate's last log entry
  - ▶ **lastLogTerm** = term of candidate's last log entry
- Response of the form (**term**, **voteGranted**)
  - ▶ **term** = current term of responding server *R*
  - ▶ **voteGranted** = *true* if *R* votes for sender; *false*, otherwise


# Server R's Response to RequestVote RPC

- **Case 1:**  $RPC.term < R.currentTerm$ 
  - R replies ( $R.currentTerm$ , false)
- **Case 2:** ( $RPC.term > R.currentTerm$ ) and **p**
  - $R.currentTerm = RPC.term$  &  $R.votedFor = RPC.candidateId$
  - R replies ( $R.currentTerm$ , true)
- **Case 3:** ( $RPC.term = R.currentTerm$ ) and ( $R.votedFor = null$ ) and **p**
  - $R.votedFor = RPC.candidateId$
  - R replies ( $R.currentTerm$ , true)
- **Case 4:** ( $RPC.term = R.currentTerm$ ) and ( $R.votedFor = RPC.candidateId$ )
  - R replies ( $R.currentTerm$ , true)
- **Case 5:** In all other cases,
  - if  $RPC.term > R.currentTerm$  then  
     $R.currentTerm = RPC.term$  &  $R.votedFor = null$
  - R replies ( $R.currentTerm$ , false)

**p:** R's log is not more complete than sender's log

# Comparing Log's Completeness

- X's log is **more complete than** Y's log if
  1. either  $X.\text{lastLogTerm} > Y.\text{lastLogTerm}$
  2. or  $(X.\text{lastLogTerm} = Y.\text{lastLogTerm})$  and  $(X.\text{lastLogIndex} > Y.\text{lastLogIndex})$



(Ongaro & Ousterhout, 2014)


# Normal Operations

- Client sends command to leader
- Leader appends command to its log
- Leader sends **AppendEntries RPCs** to all followers
- Once new log entry is **committed**
  - ▶ Leader executes command in its state machine & returns result to client
  - ▶ Leader notifies followers of committed entries in subsequent **AppendEntries RPCs**
  - ▶ Followers execute committed commands in their state machines
- **Leader Append-Only Property:** a leader never overwrites or deletes entries in its log; it only appends new entries



# Log Matching Property

- (1) If two entries in different logs have the same index and term, then they store the same command
- (2) If two entries in different logs have the same index and term, then the logs are identical in all preceding entries




(Ongaro & Ousterhout, 2014)

# Committed Log Entries


- A log entry is **directly committed** once the leader that created the entry has replicated it to a majority of servers
- All log entries preceding a directly committed entry are **indirectly committed**
- A log entry is **committed** if the entry is directly or indirectly committed

# Committed Log Entries: Example 1



- (a) S1 is leader & partially replicates new entry (2,2)
- (b) S1 fails, S5 becomes leader & appends new entry (2,3)
- (c) S5 fails, S1 becomes leader, appends new entry (3,4) & replicates entry (2,2) to S3
- (d) S1 fails, S5 becomes leader, appends new entry (3,5) & replicates entry (2,3)

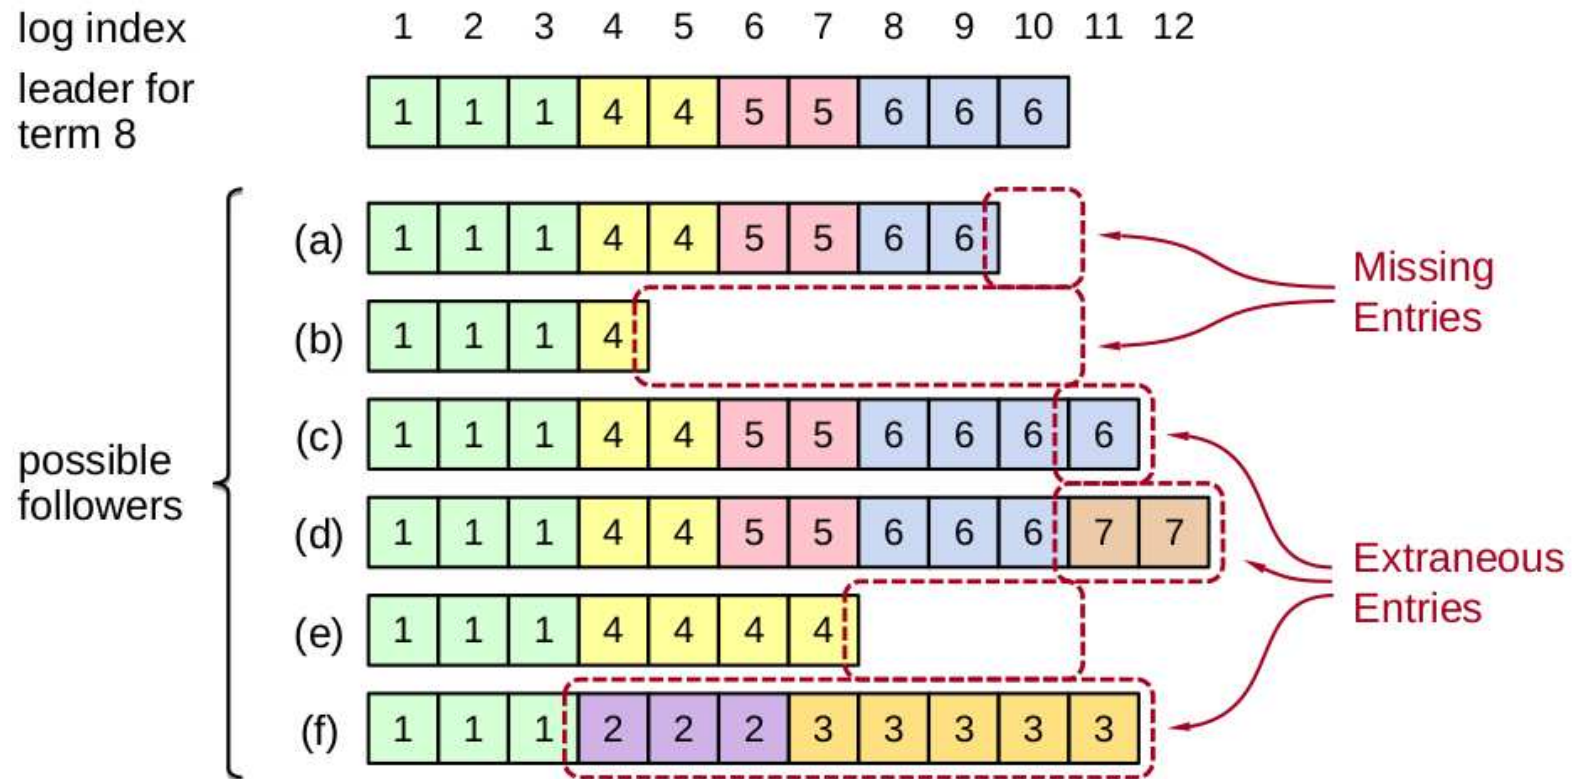
# Committed Log Entries: Example 2



- (a) S1 is leader & partially replicates new entry (2,2)
- (b) S1 fails, S5 becomes leader & appends new entry (2,3)
- (c) S5 fails, S1 becomes leader, appends new entry (3,4) & partially replicates entry (3,4)

# Log Inconsistencies


- Server failures can cause log inconsistencies



(Ongaro & Ousterhout, 2014)

# AppendEntries Consistency Check

- AppendEntries RPCs include (index,term) of entry preceding new one(s)
- Follower F must contain matching entry; otherwise F rejects AppendEntries RPC request & leader retries with lower log index



(Ongaro & Ousterhout, 2014)

# Volatile State (assumes log index starts at 1)

Volatile state on all servers (both values are initialized to 0):

- **commitIndex** - index of highest log entry known to be committed
- **lastApplied** - index of highest log entry applied to state machine

Volatile state on leaders (reinitialized after election):

- **nextIndex[]** - for each server, index of next log entry to send to that server
  - ▶ Initialized to index of leader's last log entry + 1
- **matchIndex[]** - for each server, index of highest log entry known to be replicated on server
  - ▶ Initialized to 0

# AppendEntries RPC

- AppendEntries RPC Arguments:
  - ▶ **leaderId** = identifier of leader
  - ▶ **leaderTerm** = leader's term
  - ▶ **leaderCommit** = leader's commitIndex
  - ▶ **prevLogIndex** = index of log entry immediately preceding new log entries
  - ▶ **prevLogTerm** = term of *prevLogIndex* log entry
  - ▶ **entries[]** = log entries to store
    - ★ **entries[]** is empty if AppendEntries RPC is used for heartbeat message
- Response of the form (**term**, **success**)
  - ▶ **term** = current term of responding follower F
  - ▶ **success** = *true* if F contains entry matching prevLogIndex & prevLogTerm; *false*, otherwise



# Processing AppendEntries RPC (by follower F)

1. If **leaderTerm** < F.currentTerm then reply (F.currentTerm, false)
2. If **leaderTerm** > F.currentTerm then **F.currentTerm** = **leaderTerm**
3. If **entries[]** is not empty then
  - 3.1 If F's log doesn't contain (**prevLogIndex**, **prevLogTerm**), then reply (F.currentTerm, false)
  - 3.2 If an entry e in **F's log** conflicts with a new entry (i.e., same index but different term), then delete e & all entries that follow e
  - 3.3 Append any new entries not already in **F's log**
4. If **leaderCommit** > F.commitIndex then set **F.commitIndex** = min(**leaderCommit**, index of last entry in F's log)
5. Reply (F.currentTerm, true)

# Leader Completeness Property

- **Leader Completeness Property:** if a log entry is committed in a given term, then that entry will be present in the logs of the leaders for all higher-numbered terms
- Leader election ensures that leader's log is at least as complete as a majority of servers' logs
  - ▶ This guarantees Leader Completeness Property

# Rules for all Servers

- If  $\text{commitIndex} > \text{lastApplied}$ , then
  - ▶ Increment **lastApplied** by one
  - ▶ Apply  $\log[\text{lastApplied}]$  to state machine
- If RPC request or response contains term  $T > \text{currentTerm}$ , then
  - ▶ Set **currentTerm** =  $T$
  - ▶ Convert server to follower

# Rules for Leaders

- Upon election:
  - ▶ Send initial empty AppendEntries RPC to each server
  - ▶ Repeat during idle period to prevent timeouts of election timer
- If received command from client, then
  - ▶ Append entry to **local log**
  - ▶ Respond to client after entry has been applied to state machine
- If index of last log entry  $\geq$  nextIndex for a follower F, then
  - ▶ Send AppendEntries RPC to F with log entries starting at nextIndex
  - ▶ If successful, then update **nextIndex** & **matchIndex** for F
  - ▶ Otherwise, if AppendEntries RPC fails because of log inconsistency, then decrement nextIndex & retry
- If there exists an  $N$  such that  $N > \text{commitIndex}$ , a majority of  $\text{matchIndex}[i] \geq N$ , and  $\text{log}[N].\text{term} = \text{currentTerm}$ , then
  - ▶ Set **commitIndex** =  $N$

# State Machine Safety Property

- **State Machine Safety Property:** if a server has applied a log entry at a given index to its state machine, no other server will ever apply a different log entry for the same index

# Raft topics that are not covered

- Managing cluster membership
- Log compaction
- Client interaction

# Summary of Raft Properties

- **Election Safety:** at most one leader can be elected at a given term
- **Election Liveness:** some leader must eventually be elected
- **Leader Append-Only:** a leader never overwrites or deletes entries in its log; it only appends new entries
- **Log Matching:** if two logs contain an entry with the same index and term, then the logs are identical in all entries up through the given index
- **Leader Completeness:** if a log entry is committed in a given term, then that entry will be present in the logs of the leaders for all higher-numbered terms
- **State Machine Safety:** if a server has applied a log entry at a given index to its state machine, no other server will ever apply a different log entry for the same index

# References

- D. Ongaro, J. Ousterhout, In Search of an Understandable Consensus Algorithm, USENIX Annual Technical Conference 2014
- <https://raft.github.io/>