

# **CS4224/CS5424 Lecture 10**

## **Distributed Query Optimization**

# Query Processing Steps

- **Query rewriting**
  - ▶ Query decomposition
    - ★ Translates query into relational algebra query
  - ▶ Data localization
    - ★ Rewrites distributed query into a fragment query
- **Global query optimization**
  - ▶ Finds an optimal execution plan for query
- **Distributed query execution**
  - ▶ Executes query plan to compute query result

# Why Optimize?

**Student** (sid, sname, major)

**Course** (cid, cname, area)

**Enrol** (sid, cid, grade)

```
SELECT *  
FROM Student S, Course C, Enrol E  
WHERE E.sid = S.sid  
AND E.cid = C.cid  
AND S.sid = 123
```

## Example Query Plans:

$card(R)$  denote the number of tuples in R

$card(C) = 400$


$card(E) = 10^6$

$card(S) = 10^5$


$card(\sigma_{sid=123}(S)) = 1$

$card(C \bowtie_{cid} E) = 10^6$

$card(E \bowtie_{sid} \sigma_{sid=123}(S)) = 25$




# Query Plan Search Space for $R \bowtie S \bowtie T$



# Cost Estimation of Query Plans

1. What is the evaluation cost of each operation?
  - ▶ Depends on: size of input operands, available buffer pages, available indexes, etc.
2. What is the output size of each operation?



# Cost Estimation of Query Plans (cont.)

Cost model consists of the following:

- Cost model for each operator's algorithms
- Estimation assumptions
  - ▶ **Uniformity assumption:** uniform distribution of attribute values
  - ▶ **Independence assumption:** independent distribution of values in different attributes
  - ▶ **Inclusion assumption:** inclusion dependency between join columns
- Database statistics

# Database Statistics

- $\text{length}(A)$  = size of attribute  $A$  (in bytes)
- $\text{length}(R)$  = size of tuple in relation  $R$  (in bytes)
- $\text{card}(R)$  = number of tuples in relation  $R$
- $\text{size}(R)$  = size of relation  $R$  (in bytes)
  - ▶  $\text{size}(R) = \text{card}(R) \times \text{length}(R)$
- $\text{card}(\pi_A(R))$  = number of distinct values of attribute  $R.A$
- $\min(\pi_A(R))$  = minimum value of attribute  $R.A$
- $\max(\pi_A(R))$  = maximum value of attribute  $R.A$

# Selectivity Factors

- $SF(op)$  = selectivity factor of operation  $op$ 
  - ▶ Proportion of tuples of operand relation that participate in result of operation

$$SF(\sigma_p(R)) = \frac{card(\sigma_p(R))}{card(R)}$$

$$SF(R \bowtie S) = \frac{card(R \bowtie S)}{card(R) \times card(S)}$$

$$SF(R \bowtie_A S) = \frac{card(R \bowtie_A S)}{card(R)}$$



# Selectivity Factors (cont.)

- Selectivity factors are used to estimate the cardinality of final/intermediate results
- $card(\sigma_p(R)) = SF(\sigma_p(R)) \times card(R)$
- $card(R \bowtie S) = SF(R \bowtie S) \times card(R) \times card(S)$
- $card(R \bowtie_A S) = SF(R \bowtie_A S) \times card(R)$

# Estimation of Selectivity Factors

$$SF(\sigma_{A=v}(R)) \approx \frac{1}{card(\pi_A(R))}$$

$$SF(\sigma_{A<v}(R)) \approx \frac{v - \min(\pi_A(R))}{\max(\pi_A(R)) - \min(\pi_A(R)) + 1}$$

$$SF(\sigma_{p_1 \wedge p_2}(R)) \approx SF(\sigma_{p_1}(R)) \times SF(\sigma_{p_2}(R))$$

# Estimation of Selectivity Factors

- **Inclusion assumption:** Consider  $R \bowtie_A S$   
If  $\text{card}(\pi_A(R)) \leq \text{card}(\pi_A(S))$ , then  
 $\pi_A(R) \subseteq \pi_A(S)$
- Join selectivity estimation:

$$SF(R \bowtie_A S) \approx \frac{1}{\max\{\text{card}(\pi_A(R)), \text{card}(\pi_A(S))\}}$$

- **Example:** Consider query  $Q: R \bowtie_{\text{dept}} S$

R

name	dept
Alice	CS
Bob	CS
Carol	CS
Dave	CS
Eve	CS
Fred	CS
George	CS
Henry	EE
Ivy	EE
Jane	EE

S

dept	course
CS	CS101
CS	CS111
CS	CS302
Maths	MA105
Maths	MA203
Music	MU108
Physics	PH113
Physics	PH203

$$\text{card}(R) = 10, \quad \text{card}(\pi_{\text{dept}}(R)) = 2$$

$$\text{card}(S) = 8, \quad \text{card}(\pi_{\text{dept}}(S)) = 4$$

$$\begin{aligned} \text{card}(Q) &\approx \text{card}(R) \times \frac{\text{card}(S)}{\text{card}(\pi_{\text{dept}}(S))} \\ &= 10 \times \frac{8}{4} = 20 \end{aligned}$$

# Estimation of Selectivity Factors (cont.)

$$SF(R \bowtie_A S) \approx SF_{SJ}(S.A)$$

$SF_{SJ}(S.A)$  = semijoin selectivity factor of  $S.A$

$$SF_{SJ}(S.A) = \frac{card(\pi_A(S))}{|domain(A)|}$$

# Distributed Query Optimization


- Additional complexities
  - ▶ Data fragmentation & allocation
  - ▶ Communication cost
- Optimization techniques to reduce communication cost
  - ▶ Semijoin reductions

# Search Space: Example


- **Site 1:** Course (cid, cname, area)
- **Site 2:** Enrol (sid, cid, grade)
- **Site 3:** Student (sid, sname, major)
- **Query at Site 3:**

```
SELECT  *  
FROM    Student S, Course C, Enrol E  
WHERE   E.sid = S.sid  
AND     E.cid = C.cid
```


# Search Space: Example (cont.)




Plan 1




Plan 2




Plan 3




Plan 4




Plan 5



Plan 6




Plan 7




Plan 8

..


# Query Plan Notation



$$(E \overset{1}{\bowtie} C) \overset{3}{\bowtie} S$$



$$(E \overset{2}{\bowtie} S) \overset{3}{\bowtie} C$$



$$(E \overset{2}{\bowtie} S) \overset{2}{\bowtie} C \text{ with result sent to Site 3}$$



# Optimizing Total Cost

- **Total cost** = Sum of CPU, I/O & communication costs.
- **CPU cost** =  $T_{CPU} \times$  (number of CPU instructions)
  - ▶  $T_{CPU}$  = time of a CPU instruction
- **I/O cost** =  $T_{I/O} \times$  (number of disk I/Os)
  - ▶  $T_{I/O}$  = time of a disk I/O
- **Communication cost** =  $T_{MSG} \times$  (number of messages) +  $T_{TR} \times$  (size of transferred data)
  - ▶  $T_{MSG}$  = fixed overhead for each message transmission
  - ▶  $T_{TR}$  = time to transmit one data unit

# Optimization with Semijoins

- $R \bowtie_A S = \pi_{attributes(R)}(R \Join_A S) = R \Join_A \pi_A(S)$

$$\begin{aligned} R \Join_A S &= (R \bowtie_A S) \Join_A S \\ &= (R \Join_A \pi_A(S)) \Join_A S \end{aligned}$$

- A tuple  $t \in R$  is a **dangling tuple wrt  $R \Join S$**  if  $t$  does not join with any tuple in  $S$ 
  - ▶ i.e.,  $t \notin \pi_{attributes(R)}(R \Join S)$
- $R \bowtie_A S$  eliminates dangling tuples in  $R$  (wrt  $R \Join_A S$ )

# Optimization with Semijoins: Example

Student

sid	name	major	year
1	Charlie	CS	2
2	Franklin	Maths	4
3	Lucy	Maths	3
4	Marcie	Music	2
5	Patty	Physics	4
6	Sally	CS	3

Project

pid	title	abstract	advisor	sid
1	...	...	...	5
2	...	...	...	2
3	...	...	...	3

Student  $\bowtie_{sid}$  Project

sid	name	major	year	pid	title	abstract	advisor
2	Franklin	Maths	4	2	...	...	...
3	Lucy	Maths	3	3	...	...	...
5	Patty	Physics	4	1	...	...	...

# Optimization with Semijoins: Example

Student

sid	name	major	year
1	Charlie	CS	2
2	Franklin	Maths	4
3	Lucy	Maths	3
4	Marcie	Music	2
5	Patty	Physics	4
6	Sally	CS	3

Project

pid	title	abstract	advisor	sid
1	...	...	...	5
2	...	...	...	2
3	...	...	...	3

Student  $\bowtie_{sid}$  Project


sid	name	major	year
2	Franklin	Maths	4
3	Lucy	Maths	3
5	Patty	Physics	4

(Student  $\bowtie_{sid}$  Project)  $\bowtie_{sid}$  Project


sid	name	major	year	pid	title	abstract	advisor
2	Franklin	Maths	4	2	...	...	...
3	Lucy	Maths	3	3	...	...	...
5	Patty	Physics	4	1	...	...	...

# Optimization with Semijoins

- **Example:** Site A:  $R$ , Site B:  $S$ ,  
 $size(R) < size(S)$



Direct Join Plan  
 $R \bowtie_A S$



Semijoin Plan  
 $(R \ltimes_A S) \bowtie_A S$


# Optimization with Semijoins (cont.)

- **Direct-join plan:**

- ▶ Sends  $R$  over to site of  $S$
- ▶ Joins  $R$  &  $S$  at site of  $S$

- **Semijoin plan:**

- ▶ Sends  $\pi_A(S)$  to site of  $R$
- ▶ Joins  $R$  &  $\pi_A(S)$  to eliminate dangling tuples in  $R$
- ▶ Sends non-dangling tuples of  $R$  (i.e.  $R \bowtie_A S$ ) to site of  $S$
- ▶ Joins  $R$  &  $S$  at site of  $S$




# Beneficial Semijoins

- The cost & benefit of applying semi-join optimization is defined relative to the direct-join plan
  - ▶ **Cost** measures the additional communication overhead incurred by semijoin plan over direct-join plan
  - ▶ **Benefit** measures the communication savings from semijoin plan over direct-join plan
- **Cost**( $R \ltimes_A S$ ) = cost of sending  $\pi_A(S)$
- **Benefit**( $R \ltimes_A S$ ) = savings in not sending dangling tuples of  $R$  wrt  $R \bowtie_A S$ 
  - ▶  $SF(R \ltimes_A S)$  = proportion of tuples in  $R$  that join with  $S$
  - ▶  $1 - SF(R \ltimes_A S)$  = proportion of dangling tuples in  $R$  wrt  $R \bowtie_A S$
- A semijoin is beneficial if its benefit exceeds its cost

# Beneficial Semijoins (cont.)


- Assume  $\text{size}(R) < \text{size}(S)$
- $\text{Cost}(R \bowtie_A S) = T_{MSG} + T_{TR} \times \text{size}(\pi_A(S))$
- $\text{Benefit}(R \bowtie_A S) = T_{TR} \times \text{size}(R) \times (1 - SF(R \bowtie_A S))$
- $R \bowtie_A S$  is a **beneficial semijoin** if  $\text{Benefit}(R \bowtie_A S) > \text{Cost}(R \bowtie_A S)$






# Query Plans with Semijoins

- Larger query plan search space!
- **Example:** `select * from R, S, T where R.a = S.a and R.b = T.b`



Query Plan 1



Query Plan 2

# Comparable Query Plans

- Two query plans are **comparable** if they satisfy all the following conditions:
  - (1) both plans have the same output schema,
  - (2) both plans execute their final operator on the same server, and
  - (3) the outputs of both plans are either (a) unordered or (b) sorted in the same order
- Two query plans are **incomparable** if they are not comparable
- Examples
  - ▶  $R \overset{1}{\bowtie}_p S$  and  $R \overset{1}{\ltimes}_p S$  are incomparable
  - ▶  $R \overset{1}{\bowtie}_p S$  and  $R \overset{2}{\bowtie}_p S$  are incomparable
  - ▶  $(R \overset{1}{\ltimes}_p S) \overset{2}{\bowtie}_p (S \overset{3}{\ltimes}_p R)$  and  $R \overset{2}{\bowtie}_p S$  are comparable if they satisfy condition (3)

# Query Plans

- Given a query plan  $P$ , let  $\text{cost}(P)$  denote the cost of executing  $P$
- Given two query plans,  $P1$  &  $P2$ ,  $P1$  is better than  $P2$  (or  $P2$  is worse than  $P1$ ) if (1)  $P1$  &  $P2$  are comparable and (2)  $\text{cost}(P1) < \text{cost}(P2)$
- Consider a query  $Q$  over a set of relations  $R$ , and  $S \subseteq R$ ,  $S \neq \emptyset$ . Let  $\text{optPlan}(S)$  denote the set of **optimal query subplans** of  $Q$  over  $S$ 
  - ▶ For every plan  $P \in \text{optPlan}(S)$ , there does not exist another plan  $P'$  that is better than  $P$
- Let  $\text{LOptPlan}(S)$  denote the logical query plans in  $\text{optPlan}(S)$

# Classic Dynamic Programming (DP) Algorithm

(Stocker, et al., ICDE 2001)

## Classic\_DP ( $Q$ )

Input: A SPJ query  $Q$  on relations  $R_1, \dots, R_n$

Output: A query plan for  $Q$

1. for  $i = 1$  to  $n$  do {
2.      $\text{optPlan}(\{R_i\}) = \text{accessPlans}(R_i)$
3.      $\text{prunePlans}(\text{optPlan}(\{R_i\}))$
4. }
5. for  $i = 2$  to  $n$  do
6.     for all  $S \subseteq \{R_1, \dots, R_n\}$  such that  $|S| = i$  do {
7.          $\text{optPlan}(S) = \emptyset$
8.         for all  $O \subset S$  such that  $O \neq \emptyset$  do {
9.              $\text{optPlan}(S) = \text{optPlan}(S) \cup \text{joinPlans}(\text{optPlan}(O), \text{optPlan}((S - O)))$
10.             $\text{prunePlans}(\text{optPlan}(S))$
11.         }
12.     }
13. return  $\text{optPlan}(\{R_1, \dots, R_n\})$

# Query Plan Enumeration: Example

- Distributed Database:

- ▶ Site 1: **R**(A,B,C,D)
- ▶ Site 2: **S**(X,Y)
- ▶ Site 3: **T**(E,F,G)

- Query submitted at Site 1:

```
select *  
from   R join S on R.A = S.X join T on R.D = T.F  
where R.B > 10  
and   R.C = 20  
and   T.E < 100
```

- Available indexes:  $I_B$ ,  $I_C$ ,  $I_E$
- Assumptions on database system
  - ▶ Supports only one join algorithm: hash join
  - ▶ Avoids cartesian products

# Example: Single-relation Plans

$\sigma_p(R \bowtie_{R.A=S.X} S \bowtie_{R.D=T.F} T), p = (R.B > 10) \wedge (R.C = 20) \wedge (T.E < 100)$

- **Plans for {R}**

- ▶ **Plan P1**: Table scan with “ $(B > 10) \wedge (C = 20)$ ”
- ▶ **Plan P2**: Index seek with  $I_B$  & RID-lookups with “ $C = 20$ ”
- ▶ **Plan P3**: Index seek with  $I_C$  & RID-lookups with “ $B > 10$ ”
- ▶ **Plan P4**: Index intersection with  $I_B$  &  $I_C$ , and RID-lookups
- ▶ Before pruning:  $\text{optPlan}(\{R\}) = \{P1, P2, P3, P4\}$
- ▶ Assume  $\text{cost}(P3) < \text{cost}(P4) < \text{cost}(P2) < \text{cost}(P1)$
- ▶ After pruning:  $\text{optPlan}(\{R\}) = \{P3\}$

- **Plans for {S}**

- ▶ **Plan P5**: Table scan of S
- ▶  $\text{optPlan}(\{S\}) = \{P5\}$

- **Plans for {T}**

- ▶ **Plan P6**: Table scan of T with “ $(E < 100)$ ”
- ▶ **Plan P7**: Index seek with  $I_E$  & RID-lookups
- ▶ Before pruning:  $\text{optPlan}(\{T\}) = \{P6, P7\}$
- ▶ Assume  $\text{cost}(P7) < \text{cost}(P6)$
- ▶ After pruning:  $\text{optPlan}(\{T\}) = \{P7\}$

# Example: Two-relation Plans

- Plans for  $\{R, S\}$

- Plan P8:  $\text{optPlan}(\{R\}) \overset{1}{\bowtie} \text{optPlan}(\{S\}) = P3 \overset{1}{\bowtie} P5$
- Plan P9:  $\text{optPlan}(\{R\}) \overset{2}{\bowtie} \text{optPlan}(\{S\}) = P3 \overset{2}{\bowtie} P5$
- Plan P10:  $\text{optPlan}(\{R\}) \overset{3}{\bowtie} \text{optPlan}(\{S\}) = P3 \overset{3}{\bowtie} P5$
- Plan P11:  $\text{optPlan}(\{S\}) \overset{1}{\bowtie} \text{optPlan}(\{R\}) = P5 \overset{1}{\bowtie} P3$
- Plan P12:  $\text{optPlan}(\{S\}) \overset{2}{\bowtie} \text{optPlan}(\{R\}) = P5 \overset{2}{\bowtie} P3$
- Plan P13:  $\text{optPlan}(\{S\}) \overset{3}{\bowtie} \text{optPlan}(\{R\}) = P5 \overset{3}{\bowtie} P3$
- Assume after pruning,  $\text{optPlan}(\{R, S\}) = \{P8, P9, P10\}$

- Plans for  $\{R, T\}$

- Plan P12:  $\text{optPlan}(\{R\}) \overset{1}{\bowtie} \text{optPlan}(\{T\}) = P3 \overset{1}{\bowtie} P7$
- Plan P13:  $\text{optPlan}(\{R\}) \overset{3}{\bowtie} \text{optPlan}(\{T\}) = P3 \overset{3}{\bowtie} P7$
- Plan P14:  $\text{optPlan}(\{R\}) \overset{2}{\bowtie} \text{optPlan}(\{T\}) = P3 \overset{2}{\bowtie} P7$
- Plan P15:  $\text{optPlan}(\{T\}) \overset{1}{\bowtie} \text{optPlan}(\{R\}) = P7 \overset{1}{\bowtie} P3$
- Plan P16:  $\text{optPlan}(\{T\}) \overset{3}{\bowtie} \text{optPlan}(\{R\}) = P7 \overset{1}{\bowtie} P3$
- Plan P17:  $\text{optPlan}(\{T\}) \overset{2}{\bowtie} \text{optPlan}(\{R\}) = P7 \overset{2}{\bowtie} P3$
- Assume after pruning,  $\text{optPlan}(\{R, T\}) = \{P12, P13, P17\}$

# Example: Three-relation Plans

- Plans for  $\{R, S, T\}$

- ▶ Plan P18:  $\text{optPlan}(\{R, S\}) \overset{1}{\bowtie} \text{optPlan}(\{T\})$ 
  - ★ Plan P18-8:  $P8 \overset{1}{\bowtie} P7$
  - ★ Plan P18-9:  $P9 \overset{1}{\bowtie} P7$
  - ★ Plan P18-10:  $P10 \overset{1}{\bowtie} P7$
- ▶ Plan P19:  $\text{optPlan}(\{R, S\}) \overset{2}{\bowtie} \text{optPlan}(\{T\})$
- ▶ Plan P20:  $\text{optPlan}(\{R, S\}) \overset{3}{\bowtie} \text{optPlan}(\{T\})$
- ▶ Plan P21:  $\text{optPlan}(\{T\}) \overset{1}{\bowtie} \text{optPlan}(\{R, S\})$
- ▶ Plan P22:  $\text{optPlan}(\{T\}) \overset{2}{\bowtie} \text{optPlan}(\{R, S\})$
- ▶ Plan P23:  $\text{optPlan}(\{T\}) \overset{3}{\bowtie} \text{optPlan}(\{R, S\})$
- ▶ Plan P24:  $\text{optPlan}(\{R, T\}) \overset{1}{\bowtie} \text{optPlan}(\{S\})$
- ▶ Plan P25:  $\text{optPlan}(\{R, T\}) \overset{2}{\bowtie} \text{optPlan}(\{S\})$
- ▶ Plan P26:  $\text{optPlan}(\{R, T\}) \overset{3}{\bowtie} \text{optPlan}(\{S\})$
- ▶ Plan P27:  $\text{optPlan}(\{S\}) \overset{1}{\bowtie} \text{optPlan}(\{R, T\})$
- ▶ Plan P28:  $\text{optPlan}(\{S\}) \overset{2}{\bowtie} \text{optPlan}(\{R, T\})$
- ▶ Plan P29:  $\text{optPlan}(\{S\}) \overset{3}{\bowtie} \text{optPlan}(\{R, T\})$
- ▶ Assume after pruning,  $\text{optPlan}(\{R, S, T\}) = \{P18-8, P28-17, P23-10\}$
- ▶ Optimal plan for query is the lowest cost plan in  $\text{optPlan}(\{R, S, T\})$ 
  - ★ Plans P28-17 & P23-10 require additional communication cost to send result to Site 1



# Enhanced DP Algorithm (Stocker, et al., ICDE 2001)

## Enhanced\_DP (Q)

Input: A SPJ query Q on relations  $R_1, \dots, R_n$

Output: A query plan for Q

```
1.   for i = 1 to n do {
2.       optPlan( $\{R_i\}$ ) = accessPlans( $R_i$ )
3.       prunePlans(optPlan( $\{R_i\}$ ))
4.   }
5.   for i = 2 to n do
6.       for all  $S \subseteq \{R_1, \dots, R_n\}$  such that  $|S| = i$  do {
7.           optPlan(S) =  $\emptyset$ 
8.           for all  $O \subset S$  such that  $O \neq \emptyset$  do {
N1.              for all  $P \subset O$  do {
N2.                  optPlan(S) = optPlan(S)  $\cup$  joinPlans(optPlan(O), optPlan( $(S - O) \cup P$ ), 0)
N3.                  optPlan(S) = optPlan(S)  $\cup$  SJjoinPlans(optPlan(O), optPlan( $(S - O) \cup P$ ), 0)
10.                  prunePlans(optPlan(S))
N4.              }
11.          }
N5.          timestamp = 0
N6.          do {
N7.               $\Delta$  = 'new plans with latest timestamp in S'
N8.              for all  $O \subseteq S$  such that  $O \neq \emptyset$  do {
N9.                  optPlan(S) = optPlan(S)  $\cup$  joinPlans( $\Delta$ , optPlan(O), timestamp+1)
N10.                 optPlan(S) = optPlan(S)  $\cup$  SJjoinPlans( $\Delta$ , optPlan(O), timestamp+1)
N11.                 prunePlans(optPlan(S))
N12.             }
N13.             timestamp ++
N14.         } while ( $\Delta \neq \emptyset$ )
12.     }
13.   return optPlan( $\{R_1, \dots, R_n\}$ )
```

# Incorporating Semijoins

- **Enumerator Extension**
  - ▶ Left & right operands may not be disjoint
  - ▶ Example:  $(R \bowtie S) \bowtie (T \bowtie S)$
  - ▶ Algorithm Enhanced DP: Steps N1, N2 & N3
- **Avoiding Redundant Joins & Semijoins**
  - ▶ Examples:  $R \bowtie S \bowtie S$ ,  $R \bowtie S \bowtie S$
  - ▶ Algorithm Enhanced DP: joinPlans() & SJjoinPlans()
- **Fix-point Iteration**
  - ▶ Some plans in optPlan(S) may not be complete (i.e, the plan's output schema is not from a join over S without any semijoin)
  - ▶ Example: The plan  $R \bowtie S$  in optPlan( $\{R, S\}$ ) is not complete
  - ▶ Algorithm Enhanced DP: Steps N5 to N14
- **Vertical Pruning**
  - ▶ Comparable query plans may be stored across different optPlan() entries
  - ▶ Algorithm Enhanced DP: prunePlans()


# Enumerator Extension

S	O	S-O	P	(S-O) $\cup$ P	LOptPlan(S)
$\{R, S\}$	$\{R\}$	$\{S\}$	$\{\}$	$\{S\}$	$R \bowtie S, R \ltimes S$
$\{R, S\}$	$\{S\}$	$\{R\}$	$\{\}$	$\{R\}$	$S \bowtie R, S \ltimes R$
$\{R, T\}$	$\{R\}$	$\{T\}$	$\{\}$	$\{T\}$	$R \bowtie T, R \ltimes T$
$\{R, T\}$	$\{T\}$	$\{R\}$	$\{\}$	$\{R\}$	$T \bowtie R, T \ltimes R$
$\{R, S, T\}$	$\{R\}$	$\{S, T\}$	$\{\}$	$\{S, T\}$	-
$\{R, S, T\}$	$\{S\}$	$\{R, T\}$	$\{\}$	$\{R, T\}$	$S \bowtie (R \bowtie T),$ $S \ltimes (R \bowtie T),$ $S \bowtie (R \ltimes T),$ $S \ltimes (R \ltimes T),$ $S \bowtie (T \ltimes R)$
$\{R, S, T\}$	$\{T\}$	$\{R, S\}$	$\{\}$	$\{R, S\}$	$T \bowtie (R \bowtie S),$ $T \ltimes (R \bowtie S),$ $T \bowtie (R \ltimes S),$ $T \ltimes (R \ltimes S)$
$\{R, S, T\}$	$\{R, S\}$	$\{T\}$	$\{\}$	$\{T\}$	$(R \bowtie S) \bowtie T,$ $(R \bowtie S) \ltimes T,$ $(R \ltimes S) \bowtie T,$ $(R \ltimes S) \ltimes T,$ $(S \bowtie R) \bowtie T,$ $(S \bowtie R) \ltimes T$
$\{R, S, T\}$	$\{R, S\}$	$\{T\}$	$\{R\}$	$\{R, T\}$	$(R \bowtie S) \bowtie (R \bowtie T),$ $(R \bowtie S) \bowtie (R \ltimes T),$ $(R \bowtie S) \bowtie (T \bowtie R),$ $(R \bowtie S) \bowtie (T \ltimes R),$ $(R \bowtie S) \ltimes (R \bowtie T),$ $(R \bowtie S) \ltimes (R \ltimes T),$ $(R \bowtie S) \ltimes (T \bowtie R),$ $(R \bowtie S) \ltimes (T \ltimes R),$ . . . . .


# Avoiding Redundant Joins

- Consider the example query  $Q = R \bowtie_{p_{rs}} S \bowtie_{p_{rt}} T$
- Each join predicate  $p_{rs}$  can appear in three forms in a query plan:
  - ▶ Inner-join:  $R \bowtie_{p_{rs}} S$
  - ▶ left-semijoin:  $R \ltimes_{p_{rs}} S$
  - ▶ right-semijoin:  $R \rtimes_{p_{rs}} S$
- Let **JPred(Q)** denote the set of all forms of join/semijoin predicates in a query plan for Q
- For example query Q,  $JPred(Q) = \{ \bowtie_{p_{rs}}, \ltimes_{p_{rs}}, \rtimes_{p_{rs}}, \bowtie_{p_{rt}}, \ltimes_{p_{rt}}, \rtimes_{p_{rt}} \}$
- A query plan P for a query Q is defined to be a **reasonable query plan** if for each leaf-to-root path  $L$  in P, each predicate  $j \in JPred(Q)$  appears at most once in  $L$ .
- **joinPlans()** & **SJjoinPlans()** enumerate only reasonable query plans

# Avoiding Redundant Joins (cont.)



Query Plan 1



Query Plan 2

# Fix-point Iteration

- A query plan  $P$  in  $\text{optPlan}(S)$  is **complete** if the output schema of  $P$  is the same as the output schema from a join over the relations in  $S$  (without any semijoin); otherwise it is **incomplete**
- Consider the plans obtained in  $\text{optPlan}(\{R, S\})$  at the end of step 11 in Enhanced DP Algorithm
- One of the incomplete plans in  $\text{optPlan}(\{R, S\})$  is  $R \bowtie S$
- At the end of the first fix-point iteration, the following are two complete plans that are derived from  $R \bowtie S$ :
  - ▶  $(R \bowtie S) \bowtie S$
  - ▶  $(R \bowtie S) \bowtie (S \bowtie R)$

# Vertical Pruning

- Consider a query  $Q$  over the relations  $\{R, S, T, U\}$  with join predicates  $p_{rs}$ ,  $p_{rt}$ ,  $p_{ru}$  &  $p_{st}$
- Consider the query plan  $P = R \bowtie_{p_{rs}} (S \bowtie_{p_{st}} T)$  in  $\text{optPlan}(\{R, S, T\})$
- Query plans that are comparable to  $P$  could be found in different  $\text{optPlan}()$  entries. Example:
  - ▶  $R$  in  $\text{optPlan}(\{R\})$
  - ▶  $R \bowtie_{p_{rs}} S$  in  $\text{optPlan}(\{R, S\})$
  - ▶  $R \bowtie_{p_{rt}} T$  in  $\text{optPlan}(\{R, T\})$
  - ▶  $R \bowtie_{p_{ru}} U$  in  $\text{optPlan}(\{R, U\})$
  - ▶  $R \bowtie_{p_{rt}} (T \bowtie_{p_{st}} S)$  in  $\text{optPlan}(\{R, S, T\})$
  - ▶ etc.
- **prunePlans**: Need to perform both **intra-entry pruning** as well as **inter-entry pruning**

# References

- T. Özsu & P. Valdureiz, *Distributed Query Processing*, Chapter 4, Principles of Distributed Database Systems, 4<sup>th</sup> Edition, 2020
- K. Stocker, et al., *Integrating Semi-Join-Reducers into State-of-the-Art Query Processors*, ICDE 2001