

National University of Singapore
School of Computing
CS4224/CS5424 Project
Semester 1, 2023/24

The objective of this project is to provide students the opportunity to acquire practical experience with using distributed database systems for application development. The learning tasks in this project involve the following: how to install a distributed database system on a cluster of machines, how to design a data model and implement transactions to support an application by taking into consideration of the workload of the application and the features of the database system, and how to benchmark the performance of an application.

In this project, you will be developing an application for a wholesale supplier using two different database management systems: Apache Cassandra (<http://cassandra.apache.org/>) and Citus (<https://www.citusdata.com/>).

Refer to **Canvas > Modules > Week 3 > Project Guidelines** for information on the project files to download.

1 Database Schema

The wholesale supplier has a number of geographically distributed sales districts and associated warehouses. Each regional warehouse covers 10 districts. Each district serves 3,000 customers. All warehouses maintain stocks for the 100,000 items sold by the supplier.

The relational schema for the wholesale supplier database consists of five entities (Warehouse, District, Customer, Order, Item) and two relationships (Order-Line, Stock). Order-Line specifies the items contained in each customer's order, and Stock specifies the availability information of items in each warehouse. The following subsections present the details of the schema. The key attribute(s) for each entity/relationship are shown in bold.

1.1 Warehouse

| Attribute | Meaning | Type |
|-------------|---------------------------------------|---------------|
| W_ID | Warehouse number | INT |
| W_NAME | Warehouse name | VARCHAR(10) |
| W_STREET_1 | Warehouse address | VARCHAR(20) |
| W_STREET_2 | Warehouse address | VARCHAR(20) |
| W_CITY | Warehouse address | VARCHAR(20) |
| W_STATE | Warehouse address | CHAR(2) |
| W_ZIP | Warehouse address | CHAR(9) |
| W_TAX | Warehouse sales tax rate | DECIMAL(4,4) |
| W_YTD | Year to date amount paid to warehouse | DECIMAL(12,2) |

1.2 District

| Attribute | Meaning | Type |
|---------------|--|---------------|
| D_W_ID | Warehouse number | INT |
| D_ID | District number | INT |
| D_NAME | District name | VARCHAR(10) |
| D_STREET_1 | District address | VARCHAR(20) |
| D_STREET_2 | District address | VARCHAR(20) |
| D_CITY | District address | VARCHAR(20) |
| D_STATE | District address | CHAR(2) |
| D_ZIP | District address | CHAR(9) |
| D_TAX | District sales tax rate | DECIMAL(4,4) |
| D_YTD | Year to date amount paid to district | DECIMAL(12,2) |
| D_NEXT_O_ID | Next available order number for district | INT |

D_W_ID is a foreign key that refers to Warehouse table.

1.3 Customer

| Attribute | Meaning | Type |
|----------------|---|---------------|
| C_W_ID | Warehouse number | INT |
| C_D_ID | District number | INT |
| C_ID | Customer number | INT |
| C_FIRST | Customer name | VARCHAR(16) |
| C_MIDDLE | Customer name | CHAR(2) |
| C_LAST | Customer name | VARCHAR(16) |
| C_STREET_1 | Customer address | VARCHAR(20) |
| C_STREET_2 | Customer address | VARCHAR(20) |
| C_CITY, | Customer address | VARCHAR(20) |
| C_STATE | Customer address | CHAR(2) |
| C_ZIP | Customer address | CHAR(9) |
| C_PHONE | Customer phone | CHAR(16) |
| C_SINCE | Date and time when entry was created | TIMESTAMP |
| C_CREDIT | Customer credit status | CHAR(2) |
| C_CREDIT_LIM | Customer credit limit | DECIMAL(12,2) |
| C_DISCOUNT | Customer discount rate | DECIMAL(5,4) |
| C_BALANCE | Balance of customer's outstanding payment | DECIMAL(12,2) |
| C_YTD_PAYMENT | Year to date payment by customer | FLOAT |
| C_PAYMENT_CNT | Number of payments made | INT |
| C_DELIVERY_CNT | Number of deliveries made to customer | INT |
| C_DATA | Miscellaneous data | VARCHAR(500) |

(C_W_ID, C_D_ID) is a foreign key that refers to District table.

1.4 Order

| Attribute | Meaning | Type |
|---------------|---|--------------|
| O_W_ID | Warehouse number | INT |
| O_D_ID | District number | INT |
| O_ID | Order number | INT |
| O_C_ID | Customer number | INT |
| O_CARRIER_ID | Identifier of carrier who delivered the order | INT |
| O_OL_CNT | Number of items ordered | DECIMAL(2,0) |
| O_ALL_LOCAL | Order status (whether order includes only home order-lines) | DECIMAL(1,0) |
| O_ENTRY_D | Order entry data and time | TIMESTAMP |

(O_W_ID, O_D_ID, O_C_ID) is a foreign key that refers to Customer table. The range of O_CARRIER_ID is [1,10].

1.5 Item

| Attribute | Meaning | Type |
|-------------|-----------------------|--------------|
| I_ID | Item identifier | INT |
| I_NAME | Item name | VARCHAR(24) |
| I_PRICE | Item price | DECIMAL(5,2) |
| I_IM_ID | Item image identifier | INT |
| I_DATA | Brand information | VARCHAR(50) |

1.6 Order-Line

| Attribute | Meaning | Type |
|------------------|------------------------------|--------------|
| OL_W_ID | Warehouse number | INT |
| OL_D_ID | District number | INT |
| OL_O_ID | Order number | INT |
| OL_NUMBER | Order-line number | INT |
| OL_I_ID | Item number | INT |
| OL_DELIVERY_D | Data and time of delivery | TIMESTAMP |
| OL_AMOUNT | Total price for ordered item | DECIMAL(7,2) |
| OL_SUPPLY_W_ID | Supplying warehouse number | INT |
| OL_QUANTITY | Quantity ordered | DECIMAL(2,0) |
| OL_DIST_INFO | Miscellaneous data | CHAR(24) |

(OL_W_ID, OL_D_ID, OL_O_ID) is a foreign key that refers to Order table. OL_I_ID is a foreign key that refers to Item table.

An order-line is classified as a home order-line if OL_SUPPLY_W_ID = OL_W_ID; otherwise, it is classified as a remote order-line. An order's O_ALL_LOCAL is set to *true* if and only if all its order-lines are home order-lines.

1.7 Stock

| Attribute | Meaning | Type |
|---------------|-------------------------------------|--------------|
| S_W_ID | Warehouse number | INT |
| S_I_ID | Item number | INT |
| S_QUANTITY | Quantity in stock for item | DECIMAL(4,0) |
| S_YTD | Year to date total quantity ordered | DECIMAL(8,2) |
| S_ORDER_CNT | Number of orders | INT |
| S_REMOTE_CNT | Number of remote orders | INT |
| S_DIST_01 | Information on district 1's stock | CHAR(24) |
| S_DIST_02 | Information on district 2's stock | CHAR(24) |
| S_DIST_03 | Information on district 3's stock | CHAR(24) |
| S_DIST_04 | Information on district 4's stock | CHAR(24) |
| S_DIST_05 | Information on district 5's stock | CHAR(24) |
| S_DIST_06 | Information on district 6's stock | CHAR(24) |
| S_DIST_07 | Information on district 7's stock | CHAR(24) |
| S_DIST_08 | Information on district 8's stock | CHAR(24) |
| S_DIST_09 | Information on district 9's stock | CHAR(24) |
| S_DIST_10 | Information on district 10's stock | CHAR(24) |
| S_DATA | Miscellaneous data | VARCHAR(50) |

S_I_ID is a foreign key that refers to Item table. S_W_ID is a foreign key that refers to Warehouse table.

2 Transactions

The workload for the wholesale supplier application consists of the following eight types of transactions:

1. **New Order Transaction** processes a new customer order.
2. **Payment Transaction** processes a customer payment for an order.
3. **Delivery Transaction** processes the delivery of the oldest yet-to-be-delivered order for each of the 10 districts in a specified warehouse.
4. **Order-Status Transaction** queries the status of the last order of a specified customer.
5. **Stock-Level Transaction** checks the stock level of items from a specified number of last orders at a warehouse district.
6. **Popular-Item Transaction** identifies the most popular items sold in each of a specified number of last orders at a specified warehouse district.
7. **Top-Balance Transaction** identifies the top-10 customers with the highest outstanding payment balance.
8. **Related-Customer Transaction** identifies the customers related to a specified customer.

The following subsections provide the specifications of these transaction types. The semantics of each transaction type is described in terms of a specific sequence of processing steps (you are free to implement the specification differently). For convenience, we use *null* to denote an unknown attribute value (you are free to choose an appropriate way to represent null values in your implementation). Of course, you are free to base your implementation on an entirely different schema design for each of the database system.

2.1 New Order Transaction

This transaction processes a new order from a customer.

Inputs:

1. Customer identifier (W_ID, D_ID, C_ID)
2. Number of items to be ordered NUM_ITEMS, $NUM_ITEMS \leq 20$
3. ITEM_NUMBER[i] = Item number for i^{th} item, $i \in [1, NUM_ITEMS]$
4. SUPPLIER_WAREHOUSE[i] = Supplier warehouse for i^{th} item, $i \in [1, NUM_ITEMS]$
5. QUANTITY[i] = Quantity ordered for i^{th} item, $i \in [1, NUM_ITEMS]$

Processing steps:

1. Let N denote value of the next available order number D_NEXT_O_ID for district (W_ID,D_ID)
2. Update the district (W_ID, D_ID) by incrementing D_NEXT_O_ID by one

3. Create a new order with

- O_ID = N
 - O_D_ID = D_ID
 - O_W_ID = W_ID
 - O_C_ID = C_ID
 - O_ENTRY_D = Current date and time
 - O_CARRIER_ID = *null*
 - O_OL_CNT = NUM_ITEMS
 - O_ALL_LOCAL = 0 if there exists some $i \in [1, NUM_ITEMS]$ such that SUPPLIER_WAREHOUSE[i] \neq W_ID; otherwise, O_ALL_LOCAL = 1
4. Initialize TOTAL_AMOUNT = 0

5. For i = 1 to NUM_ITEMS

- (a) Let S_QUANTITY denote the stock quantity for item ITEM_NUMBER[i] and warehouse SUPPLIER_WAREHOUSE[i]

(b) $ADJUSTED_QTY = S_QUANTITY - QUANTITY[i]$

(c) If $ADJUSTED_QTY < 10$, then set $ADJUSTED_QTY = ADJUSTED_QTY + 100$

(d) Update the stock for $(ITEM_NUMBER[i], SUPPLIER_WAREHOUSE[i])$ as follows:

- Update $S_QUANTITY$ to $ADJUSTED_QTY$
- Increment S_YTD by $QUANTITY[i]$
- Increment S_ORDER_CNT by 1
- Increment S_REMOTE_CNT by 1 if $SUPPLIER_WAREHOUSE[i] \neq W_ID$

(e) $ITEM_AMOUNT = QUANTITY[i] \times I_PRICE$, where I_PRICE is the price of $ITEM_NUMBER[i]$

(f) $TOTAL_AMOUNT = TOTAL_AMOUNT + ITEM_AMOUNT$

(g) Create a new order-line

- $OL_O_ID = N$
- $OL_D_ID = D_ID$
- $OL_W_ID = W_ID$
- $OL_NUMBER = i$
- $OL_I_ID = ITEM_NUMBER[i]$
- $OL_SUPPLY_W_ID = SUPPLIER_WAREHOUSE[i]$
- $OL_QUANTITY = QUANTITY[i]$
- $OL_AMOUNT = ITEM_AMOUNT$
- $OL_DELIVERY_D = null$
- $OL_DIST_INFO = S_DIST_xx$, where $xx = D_ID$

6. $TOTAL_AMOUNT = TOTAL_AMOUNT \times (1 + D_TAX + W_TAX) \times (1 - C_DISCOUNT)$, where W_TAX is the tax rate for warehouse W_ID , D_TAX is the tax rate for district (W_ID , D_ID), and $C_DISCOUNT$ is the discount for customer C_ID .

Output the following information:

1. Customer identifier (W_ID , D_ID , C_ID), lastname C_LAST , credit C_CREDIT , discount $C_DISCOUNT$
2. Warehouse tax rate W_TAX , District tax rate D_TAX
3. Order number O_ID , entry date O_ENTRY_D

4. Number of items NUM_ITEMS , Total amount for order $TOTAL_AMOUNT$

5. For each ordered item $ITEM_NUMBER[i]$, $i \in [1, NUM_ITEMS]$

- (a) $ITEM_NUMBER[i]$
- (b) I_NAME
- (c) $SUPPLIER_WAREHOUSE[i]$
- (d) $QUANTITY[i]$
- (e) OL_AMOUNT
- (f) $S_QUANTITY$

2.2 Payment Transaction

This transaction processes a payment made by a customer.

Inputs:

1. Customer identifier (C_W_ID, C_D_ID, C_ID)
2. Payment amount PAYMENT

Processing steps:

1. Update the warehouse C_W_ID by incrementing W_YTD by PAYMENT
2. Update the district (C_W_ID,C_D_ID) by incrementing D_YTD by PAYMENT
3. Update the customer (C_W_ID, C_D_ID, C_ID) as follows:
 - Decrement C_BALANCE by PAYMENT
 - Increment C_YTD_PAYMENT by PAYMENT
 - Increment C_PAYMENT_CNT by 1

Output the following information:

1. Customer's identifier (C_W_ID, C_D_ID, C_ID), name (C_FIRST, C_MIDDLE, C_LAST), address (C_STREET_1, C_STREET_2, C_CITY, C_STATE, C_ZIP), C_PHONE, C_SINCE, C_CREDIT, C_CREDIT_LIM, C_DISCOUNT, C_BALANCE
2. Warehouse's address (W_STREET_1, W_STREET_2, W_CITY, W_STATE, W_ZIP)
3. District's address (D_STREET_1, D_STREET_2, D_CITY, D_STATE, D_ZIP)
4. Payment amount PAYMENT

2.3 Delivery Transaction

This transaction is used to process the delivery of the oldest yet-to-be-delivered order for each of the 10 districts in a specified warehouse.

Inputs:

1. Warehouse number `W_ID`
2. Identifier of carrier assigned for the delivery `CARRIER_ID`

Processing steps:

1. For `DISTRICT_NO = 1` to 10
 - (a) Let N denote the value of the smallest order number `O_ID` for district (`W_ID`,`DISTRICT_NO`) with `O_CARRIER_ID = null`; i.e.,

$$N = \min\{t.O_ID \in \text{Order} \mid t.O_W_ID = W_ID, t.D_ID = DISTRICT_NO, t.O_CARRIER_ID = null\}$$

Let X denote the order corresponding to order number N , and let C denote the customer who placed this order

- (b) Update the order X by setting `O_CARRIER_ID` to `CARRIER_ID`
- (c) Update all the order-lines in X by setting `OL_DELIVERY_D` to the current date and time
- (d) Update customer C as follows:
 - Increment `C_BALANCE` by B , where B denote the sum of `OL_AMOUNT` for all the items placed in order X
 - Increment `C_DELIVERY_CNT` by 1

2.4 Order-Status Transaction

This transaction queries the status of the last order of a customer.

Input: Customer identifier (C_W_ID, C_D_ID, C_ID)

Output the following information:

1. Customer's name (C_FIRST, C_MIDDLE, C_LAST), balance C_BALANCE
2. For the customer's last order
 - (a) Order number O_ID
 - (b) Entry date and time O_ENTRY_D
 - (c) Carrier identifier O_CARRIER_ID
3. For each item in the customer's last order
 - (a) Item number OL_I_ID
 - (b) Supplying warehouse number OL_SUPPLY_W_ID
 - (c) Quantity ordered OL_QUANTITY
 - (d) Total price for ordered item OL_AMOUNT
 - (e) Data and time of delivery OL_DELIVERY_D

2.5 Stock-level Transaction

This transaction examines the items from the last L orders at a specified warehouse district and reports the number of those items that have a stock level below a specified threshold.

Inputs:

1. Warehouse number W_ID
2. District number D_ID
3. Stock threshold T
4. Number of last orders to be examined L

Processing steps:

1. Let N denote the value of the next available order number $D_NEXT_O_ID$ for district (W_ID, D_ID)
2. Let S denote the set of items from the last L orders for district (W_ID, D_ID) ; i.e.,

$$S = \{t.OL_I_ID \mid t \in \text{Order-Line}, t.OL_D_ID = D_ID, t.OL_W_ID = W_ID, t.OL_O_ID \in [N-L, N)\}$$

3. Output the total number of items in S where its stock quantity at W_ID is below the threshold;
i.e., $S_QUANTITY < T$

2.6 Popular-Item Transaction

This transaction finds the most popular item(s) in each of the last L orders at a specified warehouse district. Given two items X and Y in the same order O , X is defined to be more popular than Y in O if the quantity ordered for X in O is greater than the quantity ordered for Y in O .

Inputs:

1. Warehouse number W_ID
2. District number D_ID
3. Number of last orders to be examined L

Processing steps:

1. Let N denote the value of the next available order number $D_NEXT_O_ID$ for district (W_ID, D_ID)
2. Let S denote the set of last L orders for district (W_ID, D_ID) ; i.e.,

$$S = \{t.O_ID \mid t \in \text{Order}, t.O_D_ID = D_ID, t.O_W_ID = W_ID, t.O_ID \in [N - L, N)\}$$

3. For each order number x in S

(a) Let I_x denote the set of order-lines for this order; i.e.,

$$I_x = \{t \in \text{Order-Line} \mid t.OL_O_ID = x, t.OL_D_ID = D_ID, t.OL_W_ID = W_ID\}$$

(b) Let $P_x \subseteq I_x$ denote the subset of popular items in I_x ; i.e.,

$$t \in P_x \iff \forall t' \in I_x, t'.OL_QUANTITY \leq t.OL_QUANTITY$$

Output the following information:

1. District identifier (W_ID, D_ID)
2. Number of last orders to be examined L
3. For each order number x in S
 - (a) Order number O_ID & entry date and time O_ENTRY_D
 - (b) Name of customer who placed this order $(C_FIRST, C_MIDDLE, C_LAST)$,
 - (c) For each popular item in P_x
 - i. Item name I_NAME
 - ii. Quantity ordered $OL_QUANTITY$
4. The percentage of examined orders that contain each popular item
 - (a) For each distinct popular item $t \in P_x, x \in S$
 - i. Item name I_NAME
 - ii. The percentage of orders in S that contain the popular item t

2.7 Top-Balance Transaction

This transaction finds the top-10 customers ranked in descending order of their outstanding balance payments.

Processing steps:

1. Let $C \subseteq \text{Customer}$ denote the subset of 10 customers (i.e., $|C| = 10$) such that for each pair of customers (x, y) , where $x \in C$ and $y \in \text{Customer} - C$, we have $x.C_BALANCE \geq y.C_BALANCE$.

Output the following information:

1. For each customer in C ranked in descending order of $C_BALANCE$:
 - (a) Name of customer (C_FIRST , C_MIDDLE , C_LAST)
 - (b) Balance of customer's outstanding payment $C_BALANCE$
 - (c) Warehouse name of customer W_NAME
 - (d) District name of customer D_NAME

2.8 Related-Customer Transaction

This transaction finds all the customers who are related to a specific customer. Given a customer C , another customer C' is defined to be related to C if all the following conditions hold:

- C and C' are associated with different warehouses; i.e., $C.C_W_ID \neq C'.C_W_ID$, and
- C and C' each has placed some order, O and O' , respectively, where both O and O' contain at least two items in common.

Input: Customer identifier (C_W_ID, C_D_ID, C_ID)

Processing steps:

1. Let S be the set of customers who are related to the customer identified by (C_W_ID, C_D_ID, C_ID).

$$\begin{aligned}
 S = \{ & C' \in Customer \mid C'.C_W_ID \neq C_W_ID, \\
 & \exists O \in Order, O.O_W_ID = C_W_ID, O.O_D_ID = C_D_ID, O.O_C_ID = C_ID, \\
 & \exists O' \in Order, O'.O_W_ID = C'.C_W_ID, O'.O_D_ID = C'.C_D_ID, O'.O_C_ID = C'.C_ID \\
 & \exists OL1 \in OrderItem, OL1.OL_W_ID = O.O_W_ID, OL1.OL_D_ID = O.O_D_ID, OL1.OL_O_ID = O.O_ID, \\
 & \exists OL2 \in OrderItem, OL2.OL_W_ID = O.O_W_ID, OL2.OL_D_ID = O.O_D_ID, OL2.OL_O_ID = O.O_ID, \\
 & \exists OL1' \in OrderItem, OL1'.OL_W_ID = O'.O_W_ID, OL1'.OL_D_ID = O'.O_D_ID, OL1'.OL_O_ID = O'.O_ID, \\
 & \exists OL2' \in OrderItem, OL2'.OL_W_ID = O'.O_W_ID, OL2'.OL_D_ID = O'.O_D_ID, OL2'.OL_O_ID = O'.O_ID, \\
 & OL1.OL_I_ID \neq OL2.OL_I_ID, OL1'.OL_I_ID \neq OL2'.OL_I_ID, \\
 & OL1.OL_I_ID = OL1'.OL_I_ID, OL2.OL_I_ID = OL2'.OL_I_ID \}
 \end{aligned}$$

Output the following information:

1. Customer identifier (C_W_ID, C_D_ID, C_ID)
2. For each customer $C \in S$
 - (a) Output the identifier of C

2.9 Transaction Workload

In this project, we consider a transaction workload with the following distribution for the eight transaction types:

| Transaction Type | Frequency (%) |
|------------------|---------------|
| New-Order | 40 |
| Payment | 20 |
| Delivery | 20 |
| Order-Status | 4 |
| Stock-Level | 4 |
| Popular-Item | 4 |
| Top-Balance | 6 |
| Related-Customer | 2 |

The transactions satisfy the following properties:

1. Number of items per new order: $[4, 20]$
2. Stock threshold T in Stock-level transactions: $[1, 10]$
3. Number L in Stock-level and Popular-item transactions: $[20, 50]$

3 What to do

The project consists of two parts with each part comprising of four main tasks. The first part is to be done using [Apache Cassandra 4.1.3](#), and the second part is to be done using [Citrus 12.0 \(with PostgreSQL version 15.4\)](#). The project deliverables are described in Section 4. The maximum score for each part is 15 marks.

For each part (using a specific database system) the goal is to make the best use of the features provided by the database system to design and develop an efficient approach to process the transaction workload. Specifically, for each database system, there are four tasks:

1. **Data Modeling.** The first task is to design a data model to support the application's transaction workload.
2. **Transaction Implementation.** The second task is to implement the application with the following nine functions/programs:
 - (a) A program/script for loading data into your database.
 - (b) One function for each of the eight transaction types.
 - (c) **A main driver program that simulates a client executing transactions on the database.** This driver function reads its inputs from `stdin`, where each input specifies an instance of one of the eight transaction types. For each transaction read, the driver invokes the appropriate transaction function to execute that transaction. The driver repeatedly reads a transaction from `stdin` and processes that transaction until EOF is read. The driver function must process the transactions sequentially; i.e., the current transaction must be completed before reading and processing the next transaction. Any output to be produced by a transaction should be written into `stdout`. At the end of processing all the transactions, the driver outputs to `stderr` the following information:
 - Total number of transactions processed
 - Total elapsed time for processing the transactions (in seconds)
 - Transaction throughput (number of transactions processed per second)
 - Average transaction latency (in ms)
 - Median transaction latency (in ms)
 - 95th percentile transaction latency (in ms)
 - 99th percentile transaction latency (in ms)

The input specification for each transaction type consists of at least one line with each line consisting of comma-separated parameter values. The first value in the first line must be either N , P , D , O , S , I , T , or R , denoting, respectively, a new order, payment, delivery, order-status, stock-level, popular-item, top-balance, or related-customer transaction. The input specification for each of the transaction types is defined as follows:

- (a) **New Order Transaction** consists of $M+1$ lines, where M denote the number of items in the new order. The first line consists of five comma-separated values: N, C_ID, W_ID, D_ID, M .

Each of the M remaining lines specifies an item in the order and consists of three comma-separated values: `OLI.ID,OL.SUPPLY_W_ID,OL.QUANTITY`.

- (b) **Payment Transaction** consists of one line of input with five comma-separated values: `P,C_W_ID,C_D_ID,C_ID,PAYMENT`.
- (c) **Delivery Transaction** consists of one line of input with three comma-separated values: `D,W_ID,CARRIER_ID`.
- (d) **Order-Status Transaction** consists of one line of input with four comma-separated values: `O,C_W_ID,C_D_ID,C_ID`.
- (e) **Stock-Level Transaction** consists of one line of input with five comma-separated values: `S,W_ID,D_ID,T,L`.
- (f) **Popular-Item Transaction** consists of one line of input with four comma-separated values: `I,W_ID,D_ID,L`.
- (g) **Top-Balance Transaction** consists of one line of input with one value: `T`.
- (h) **Related-Customer Transaction** consists of one line of input with four comma-separated values: `R,C_W_ID,C_D_ID,C_ID`.

3. Installation.

The third task is to install and configure the database system on a cluster of 5 servers. You will be using SoC's Compute Cluster for this project; refer to Appendix A for details on the cluster usage. Your database should be installed in each server's local disk at `/temp` (not to be confused with `/tmp`). An initial set of data will be provided for you to initialize your database with data.

For Cassandra, the database is to be partitioned among 5 servers using a replication factor of 3 (i.e., there are three copies of each database record).

For Citus, the database should be configured with one coordinator node and four worker nodes. Each distributed table is to be partitioned among the four worker nodes.

4. **Performance Measurement.** The fourth task is to measure the performance of your application with 20 clients (C_0, C_1, \dots, C_{19}) executing transactions concurrently from the cluster of 5 servers (S_0, S_1, S_2, S_3, S_4) using a set of 20 transaction specification files (`0.txt, 1.txt, \dots, 19.txt`). The transaction specification files will be provided to you. Each client C_i is simulated by executing an instance of your main driver program on server $S_{(i \bmod 5)}$ which reads from the transaction specification file named `i.txt`. Thus, there are 4 clients executing transactions concurrently on each server.

Report the following performance measurements:

- (1) As stated in task 2(c), on completion of its transaction executions, each client reports the following performance measurements (with all floating point measurement values rounded off to two decimal places):
 - (a) Number of executed transactions
 - (b) Total transaction execution time (in seconds)

- (c) Transaction throughput (i.e., number of executed transactions per second given by (a) / (b))
- (d) Average transaction latency (in ms)
- (e) Median transaction latency (in ms)
- (f) 95th percentile transaction latency (in ms)
- (g) 99th percentile transaction latency (in ms)

Record these 7 measurements in a single 8-column CSV file named **clients.csv**, with the following schema for each line:

client_number, measurement_a, ..., measurement_g

This file should have exactly 20 lines.

- (2) The minimum, average, and maximum transaction throughputs among the 20 clients. Record these 3 measurements in a single 3-column CSV file named **throughput.csv**, with the following schema for each line:

min_throughput, max_throughput, avg_throughput

- (3) Report the database state at the end of the performance measurement by computing the following 15 statistics from the database; for convenience, we specify the semantics of these 15 statistics using the following 6 SQL queries:

- i. select sum(W_YTD) from Warehouse
- ii. select sum(D_YTD), sum(D_NEXT_O_ID) from District
- iii. select sum(C_BALANCE), sum(C_YTD_PAYMENT), sum(C_PAYMENT_CNT), sum(C_DELIVERY_CNT) from Customer
- iv. select max(O_ID), sum(O_OL_CNT) from Order
- v. select sum(OL_AMOUNT), sum(OL_QUANTITY) from Order-Line
- vi. select sum(S_QUANTITY), sum(S_YTD), sum(S_ORDER_CNT), sum(S_REMOTE_CNT) from Stock

Record these 15 statistics in a 15-line file named **dbstate.csv**, with one statistic on one line:

| |
|---------------------|
| sum(W_YTD) |
| sum(D_YTD) |
| sum(D_NEXT_O_ID) |
| sum(C_BALANCE) |
| sum(C_YTD_PAYMENT) |
| sum(C_PAYMENT_CNT) |
| sum(C_DELIVERY_CNT) |
| max(O_ID) |
| sum(O_OL_CNT) |
| sum(OL_AMOUNT) |
| sum(OL_QUANTITY) |
| sum(S_QUANTITY) |
| sum(S_YTD) |
| sum(S_ORDER_CNT) |
| sum(S_REMOTE_CNT) |

As the performance of these statistic computations will not be measured, your database schema design should be optimized for the performance of only the transaction executions (and not for these statistics computations).

3.1 Optional Task (Maximum of 2 bonus marks)

If you have completed the above four tasks for each of Cassandra and Citus, you are encouraged to conduct additional performance experiments by comparing against other database configurations. For example, you may want to explore using a different database data model or explore a different implementation based on using different database features. As another example, you may want to explore the performance of Cassandra without any data replication.

4 Project Deliverables

The project consists of two deliverables with the following submission deadlines:

| Deliverable | Deadline |
|----------------------------------|------------------------------|
| One-page progress report | October 6 (Friday), 23:59 |
| Project report & code submission | November 9 (Thursday), 23:59 |

4.1 Progress Report

1. A progress report (must be in **pdf format**, maximum of 1 page, minimum font size 10). The progress report should include the following:
 - (a) Team identifier and names of all team members.
 - (b) For each team member, list of his/her allocated tasks.
 - (c) Progress for Cassandra implementation
 - i. A list of completed tasks.
 - ii. A list of in-progress tasks.
 - (d) Progress for Citus implementation
 - i. A list of completed tasks.
 - ii. A list of in-progress tasks.
 - (e) Any other comments (optional).

4.2 Final Report & Code

1. A report (must be in **pdf format**, maximum of 30 pages, minimum font size 10). The report should include the following:
 - (a) Team identifier and names of all team members on the first page.
 - (b) For each team member, state his/her contributions to the project.

- (c) For each part (Cassandra and Citus implementation)
 - i. Description of the data models with justifications for the design decisions.
 - ii. A brief description of your implementation of the transaction functions. You should not repeat the contents of the transaction specifications here; instead, focus on explaining why your implementation is efficient, etc.
 - iii. A brief discussion of the essential configuration options used
 - iv. A discussion of the results for your performance benchmarking.
 - (d) A discussion of any lessons learnt, difficulties encountered, etc.
 - (e) A list of all references (e.g., books, papers, websites) consulted for the project.
2. For each part (Cassandra and Citus implementation)
- (a) Source code of your application (including all scripts/configuration files). Marks will be penalized for incomplete submissions.
 - (b) A detailed README file describing how to configure, setup and run your application using your submitted code/configuration files.
 - (c) The files **clients.csv**, **throughput.csv**, and **dbstate.csv**.

4.3 How to Submit

Each team should submit four files in total; do not make multiple submissions of each file.

Upload your project deliverables into the appropriate **Canvas** > **Assignments** folders as follows:

- **Project progress report:** Name your file as **X_progress.pdf**
- **Project final report:** Name your file as **X_final.pdf**
- **Cassandra source code:** Name your file as **X_cassandra.zip**
- **Citus source code:** Name your file as **X_citus.zip**

In the above file names, "X" denote your team identifier (i.e., a, \dots, z).

Late submission policy: Any submission that is late by one day will incur a penalty of 2 marks; any submission that is late by two days will incur a penalty of 4 marks; and any submission that is late by more than two days will not be graded and will receive 0 marks. Late days are rounded up to the nearest integer. For example, a submission that is late by 1 hour will count as 1 day late. So if your Cassandra code submission is late by 5 hours and your final report is late by 30 hours, your total penalty will be 6 marks.

5 Project Evaluation

1. It is important that your submitted README file clearly explains how to configure and run your implementation, and how to compile and generate any required executable files.

2. It is important that your implementation can be executed using the command line interface without the need to install any IDE environment.
3. We will not be making any modifications to your source code. It is therefore important that all parameter configurations (e.g., IP address) be specified in scripts instead of in source code.
4. Your project will be graded based on the following criteria: data model design, implementation of transactions, experimental evaluations, and written report.
5. There will a 20-minute project evaluation session for each team during week 13. Each evaluation session involves a combination of Q&A and demo of implemented code. All team members are expected to attend the evaluation session. The registration of the evaluation sessions will be announced in week 12.

6 References

The following are some useful references for the project.

6.1 Apache Cassandra

1. Apache Cassandra 4.1 <https://cassandra.apache.org/doc/4.1/index.html>
2. DataStax documentation for Apache Cassandra <https://docs.datastax.com/en/home/docs/cassandra.html>
3. Jeff Carpenter & Eben Hewitt, Cassandra: The Definitive Guide: Distributed Data at Web Scale, Revised 3rd Edition, O'Reilly Media, 2022
<https://www.oreilly.com/library/view/cassandra-the-definitive/9781492097136/>
4. Jeff Carpenter & Eben Hewitt, Cassandra: The Definitive Guide: Distributed Data at Web Scale, 2nd Edition, O'Reilly Media, 2016 <http://linc.nus.edu.sg/search/Y?SEARCH=9781491933633>

6.2 Citus

1. <https://www.citusdata.com/>
2. <https://docs.citusdata.com/en/v12.0>

6.3 PostgreSQL

1. <https://www.postgresql.org/>
2. <https://www.postgresql.org/docs/15/index.html>